


Systems Programming

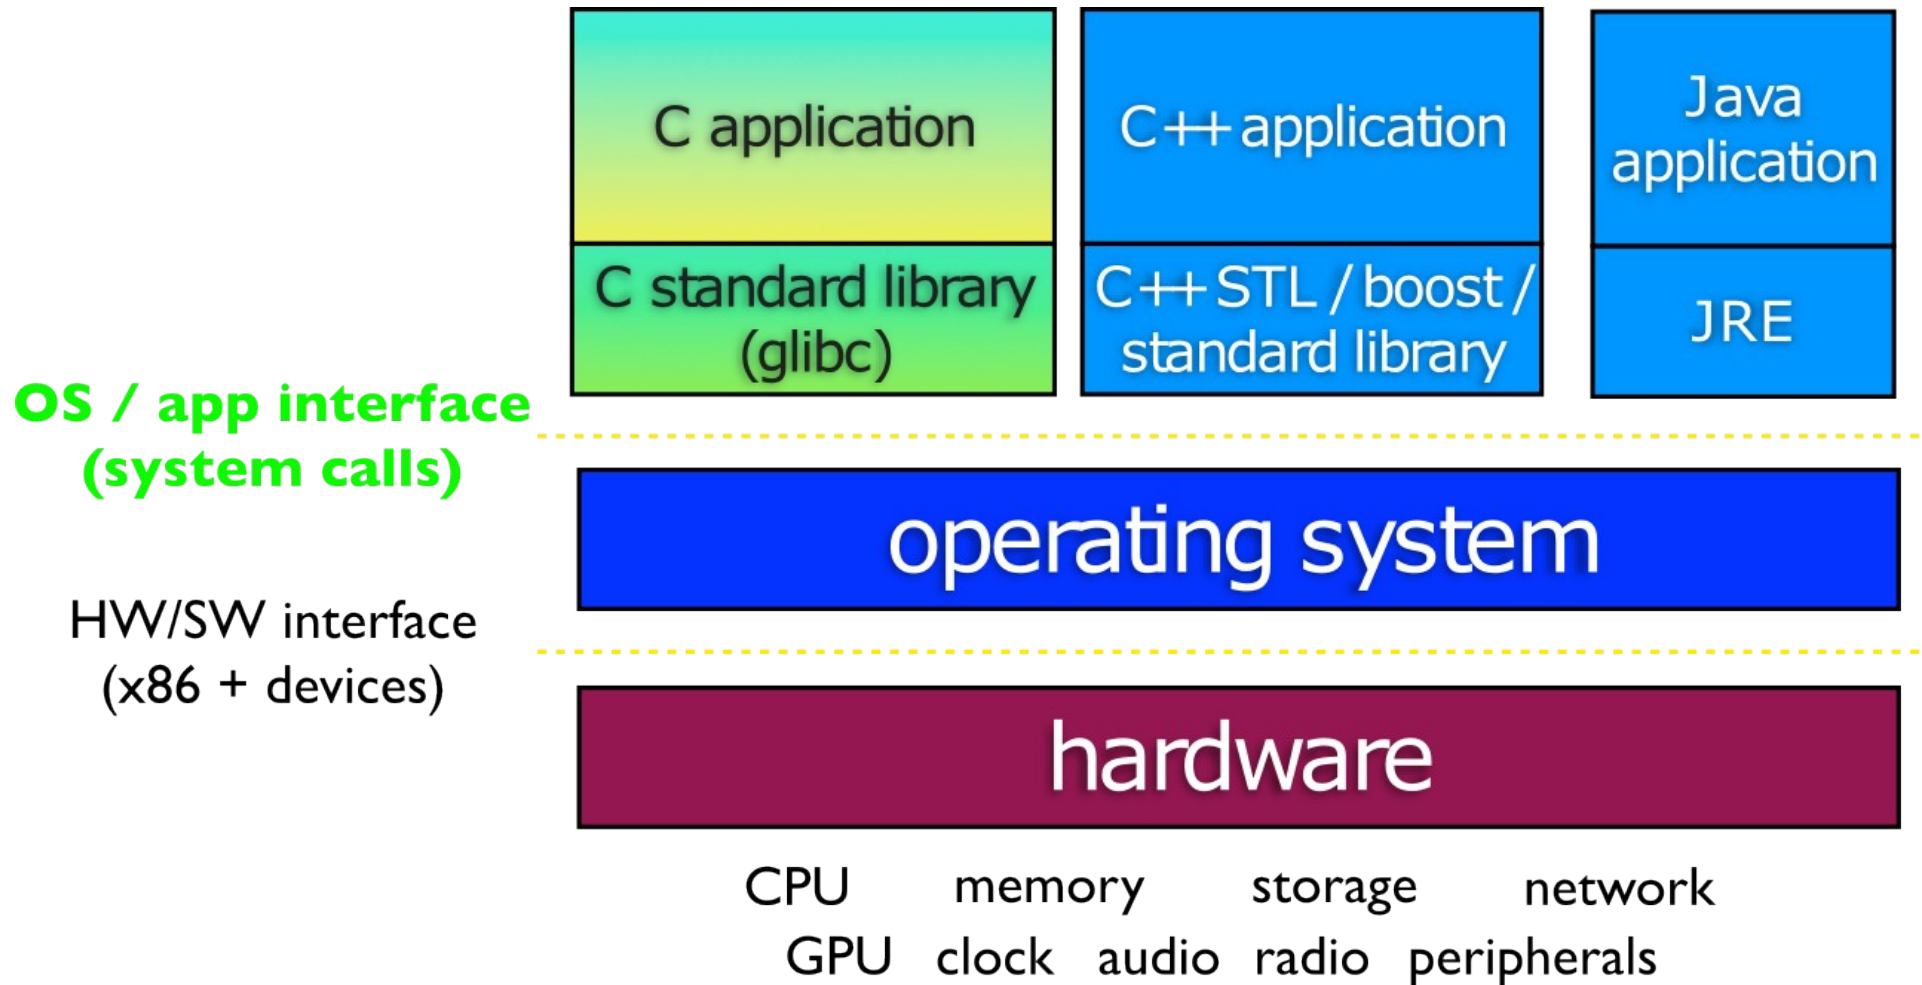
Devin J. Pohly <djpohly@cse.psu.edu>

Software system

- A platform, application, or other structure that:
 - is composed of multiple modules
 - the system's **architecture** defines the *interfaces of* and *relationships between* the modules
 - usually is complex in terms of implementation, performance, and management
 - hopefully meets some requirements
 - performance, security, fault tolerance, data consistency




10,000-foot view




A layered view

- Each layer:
 - provides service to layers above
 - understands and relies on layers below



A layered view

- Higher-level layers
 - more useful, portable, reliable abstractions
- Lower-level layers
 - constrained by performance, footprint, behavior of the layers below




Example system

- Operating system
 - a software layer that abstracts away the messy details of hardware into a useful, portable, powerful interface
 - modules:
 - filesystem, virtual memory management, network stack, protection system, scheduler
 - each of these “subsystems” is a major system of its own!
 - design and implementation has many engineering tradeoffs
 - e.g., speed vs (portability, maintainability, simplicity)



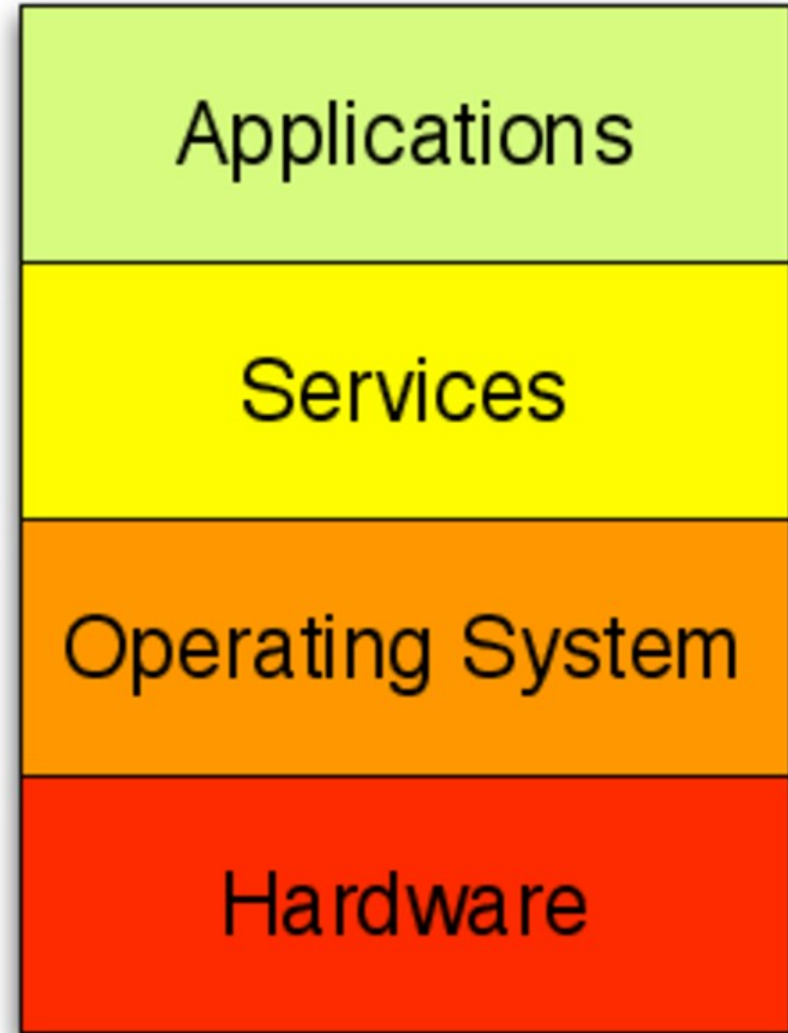
Another example system

- Web server framework
 - a software layer that abstracts away the messy details of OSes, HTTP protocols, and storage systems to simplify building powerful, scalable Web services
 - modules such as:
 - HTTP server, HTML template system, database storage, user authentication
 - also has many, many tradeoffs:
 - programmer convenience vs. performance
 - simplicity vs. extensibility
- Note: we will focus on the OS as an example system



Systems and layers

- Layers are collections of system functions that support some *abstraction* to service/app above
 - Hides the specifics of the implementation of the layer
 - Hides the specifics of the layers below
 - Abstraction may be provided by software or hardware
 - Examples from the OS: processes, virtual memory, files



A real-world abstraction



What does this do?

A real-world abstraction




How about this?

(Side note: “Law of Leaky Abstractions”)

Processes

- Processes are independent programs running concurrently within the operating system
 - The execution abstraction provides the illusion that each process has sole control of the entire computer (a single stack and execution context)
- PRO TIP: if you want to see what processes are running on a UNIX system, use the `ps` command. Try “`ps -ax`”.



Virtual memory

- The *virtual memory* abstraction provides control over an imaginary address space
 - Each process has its own virtual address space
 - The OS/hardware work together to map the address onto:
 - Physical memory addresses
 - Addresses on disk (*swap space*)



Virtual memory

- Advantages of virtual memory
 - Allows process to use entire address space
 - Avoids interference from other processes
 - Swap allows more memory use than physically available



- A file is an abstraction of a read-only, write-only, or read/write data object.
- A *data file* is a collection of data on some medium
 - often on secondary storage (hard disk)
 - also called a “regular file”
- What other “objects” could fit this abstraction?



- In UNIX nearly *everything* is a file
 - Devices like printers, USB buses, disks, etc.
 - System services like sources of randomness (RNG)
 - Terminals (user I/O devices)
 - Even process information!
- PROTIP: The `/dev` directory of UNIX contains real and virtual devices. Try “`ls /dev`”.



Systems programming

- The tools you need to build a system using these abstractions
 - **programming skills:** C (the abstraction for ISA)
 - **engineering discipline:** testing, debugging, performance analysis
 - **knowledge:** long list of interesting topics
 - concurrency, OS interfaces and semantics, techniques for consistent data management, algorithms, distributed systems
 - most important: deep understanding of the “layer below”



Programming languages

- *Assembly language* (ASM) and machine language
 - (approximately) directly executed by hardware
 - tied to a specific machine architecture, not portable
 - no notion of structure, few programmer conveniences
 - possible to write really, really fast code

IDA View-B

```
.text:00425690  
.text:00425690 ; !!!!!!!!!!!!!!! S U B R O U T I N E !!!!!!!!!!!!!!!  
.text:00425690 ; Attributes: library function bp-based frame  
.text:00425690 ; char *__cdecl strdup(const char *s)  
.text:00425690 _strdup proc near  
.text:00425690 s = dword ptr 8  
.text:00425690  
.text:00425690 push ebp  
.text:00425691 mov ebp, esp  
.text:00425693 push ebx  
.text:00425694 push esi  
.text:00425695 push edi  
.text:00425696 mov edi, [ebp+s]  
.text:00425699 push edi  
.text:0042569A call _strlen  
.text:0042569F pop ecx  
.text:004256A0 mov esi, eax  
.text:004256A2 inc esi  
.text:004256A3 push esi  
.text:004256A4 call _malloc  
.text:004256A9 pop ecx  
.text:004256AA mov ebx, eax  
.text:004256AC test eax, eax  
.text:004256AE jz short end  
.text:004256B0 push esi  
.text:004256B1 push edi  
.text:004256B2 push ebx  
.text:004256B3 call _memcpy  
.text:004256B8 add esp, 0Ch  
.text:004256B8 end:  
.text:004256B8 mov eax, ebx  
.text:004256BD pop edi
```

Function calls: _strdup

Address	Caller	Instruction
.text:004117F0	sub_4116BC	call _strdup
.text:00411823	sub_4116BC	call _strdup
.text:00423206	__setLocale32A	call _strdup
.text:0042A62E	__setMonetary	call _strdup
.text:0042A439	__setTime	call _strdup
.text:0042A462	__setTime	call _strdup
.text:0042A48B	__set	call _strdup
.text:0042D815	__set	call _strdup


Address Called function

.text:0042569A	call _strlen
.text:004256A4	call _malloc
.text:004256B3	call _memcpy

```
add esp, 0Ch  
mov ecx, [ebx+8]  
push ecx ; block  
call _free  
pop ecx  
lea eax, [ebp+s]  
push eax ; s  
call _strdup  
pop ecx  
mov [ebx+8], eax  
push 40h ; n  
lea edx, [ebp+s]  
push edx ; s  
mov ecx, [ebx+0Ch]  
push ecx ; int  
call __win32DateTimeToPOSIX  
add esp, 0Ch  
mov eax, [ebx+0Ch]  
push eax ; block  
call _free
```


Programming languages

- *Compilation* of a programming language results in executable code to be run by hardware.
- gcc (C compiler) produces target machine executable code (ISA)
- javac (Java compiler) produces Java Virtual Machine executable code



Programming languages

- Structured but low-level languages (C and C++)
 - hide some architectural details
 - kind of portable
 - have a few useful abstractions like types, arrays, procedures, objects



Programming languages

- C permits (or forces?) the programmer to handle low-level details like memory management, locks, threads
- Low-level enough to be **fast** and to give the programmer **control** over resources
 - double-edged sword: low-level enough to be complex, error-prone
 - shield: engineering discipline



Programming languages

- High-level languages (Python, Ruby, JavaScript, ...)
 - focus on productivity and usability over performance
 - powerful abstractions to hide the low-level gritty details (bounded arrays, garbage collection, rich libraries)
 - usually interpreted, translated, or compiled via an intermediate representation
 - slower (by 1.2x-10x), less control



-

- Tools
 - gcc, gdb, g++, objdump, nm, gcov/lcov, valgrind, IDEs, race detectors, model checkers
- Lower-level systems
 - UNIX system call API, relational databases, map/reduce, Django
- Systems foundations
 - transactions, two-phase commit, consensus, RPC, virtualization, cache coherence, applied crypto