



Unix and Operating Systems (Part 2)

Devin J. Pohly <djpohly@cse.psu.edu>

Introduction to Systems...

- Crossover knowledge from CMPEN 331 and CMPSC 473.
- Last day of dry background material for now.
- The programming and cool Unix stuff is coming, I promise!



Types of storage

- Example: HDD vs. SSD
 - What practical differences are there?
- How about RAM?
- Google Drive?



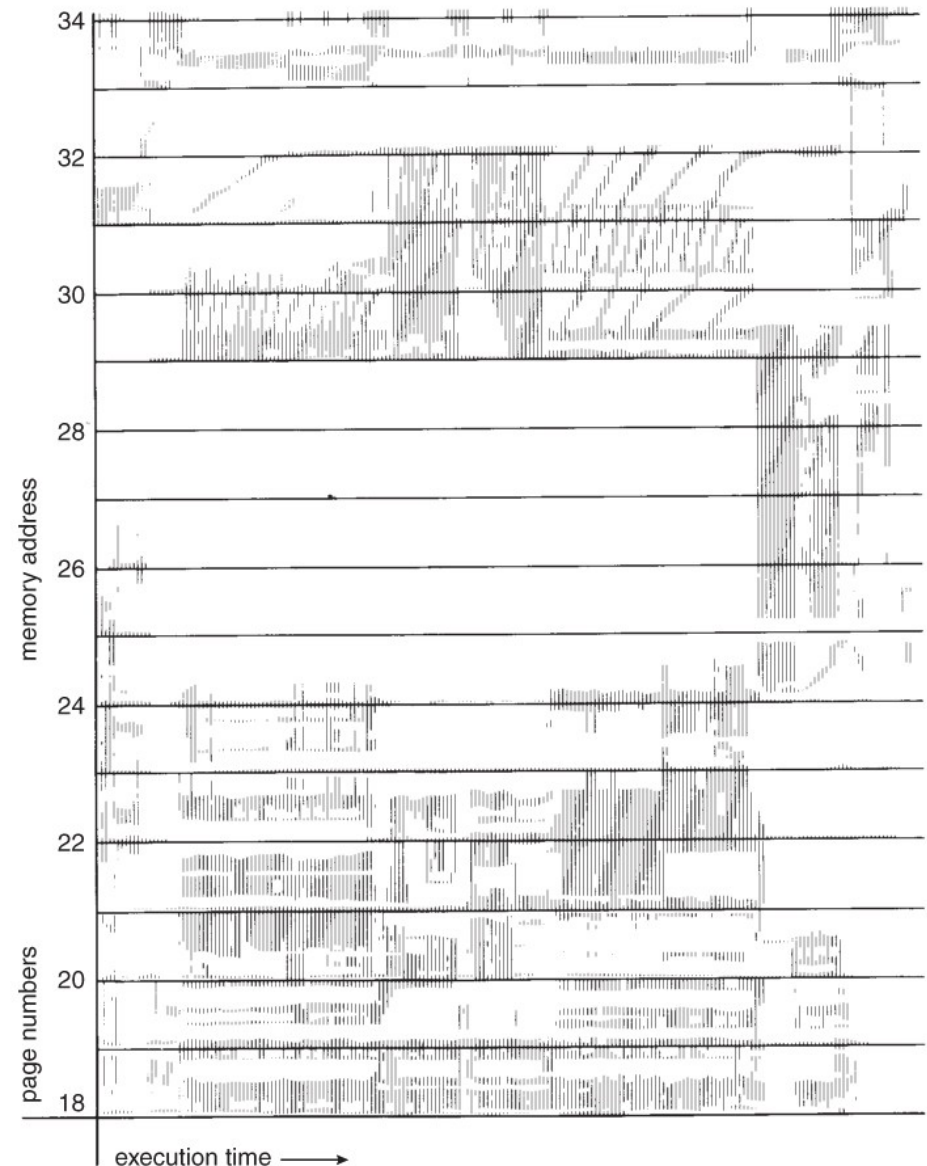
A few observations

- Some fundamental and/or enduring properties of hardware and software:
 - Fast storage technologies cost more per byte, have less capacity, and require more power (hotter!).
 - The gap between CPU and main memory speed is widening.
 - And something called *locality*...



Locality

- Temporal locality
 - Memory recently referenced is likely to be referenced again soon...
- Spatial locality
 - ... and so is the memory around it.
- Well-written programs tend to exhibit good locality.



Memory hierarchies

- These fundamental properties complement each other beautifully.
- They suggest an approach for organizing memory and storage systems known as a **memory hierarchy**.

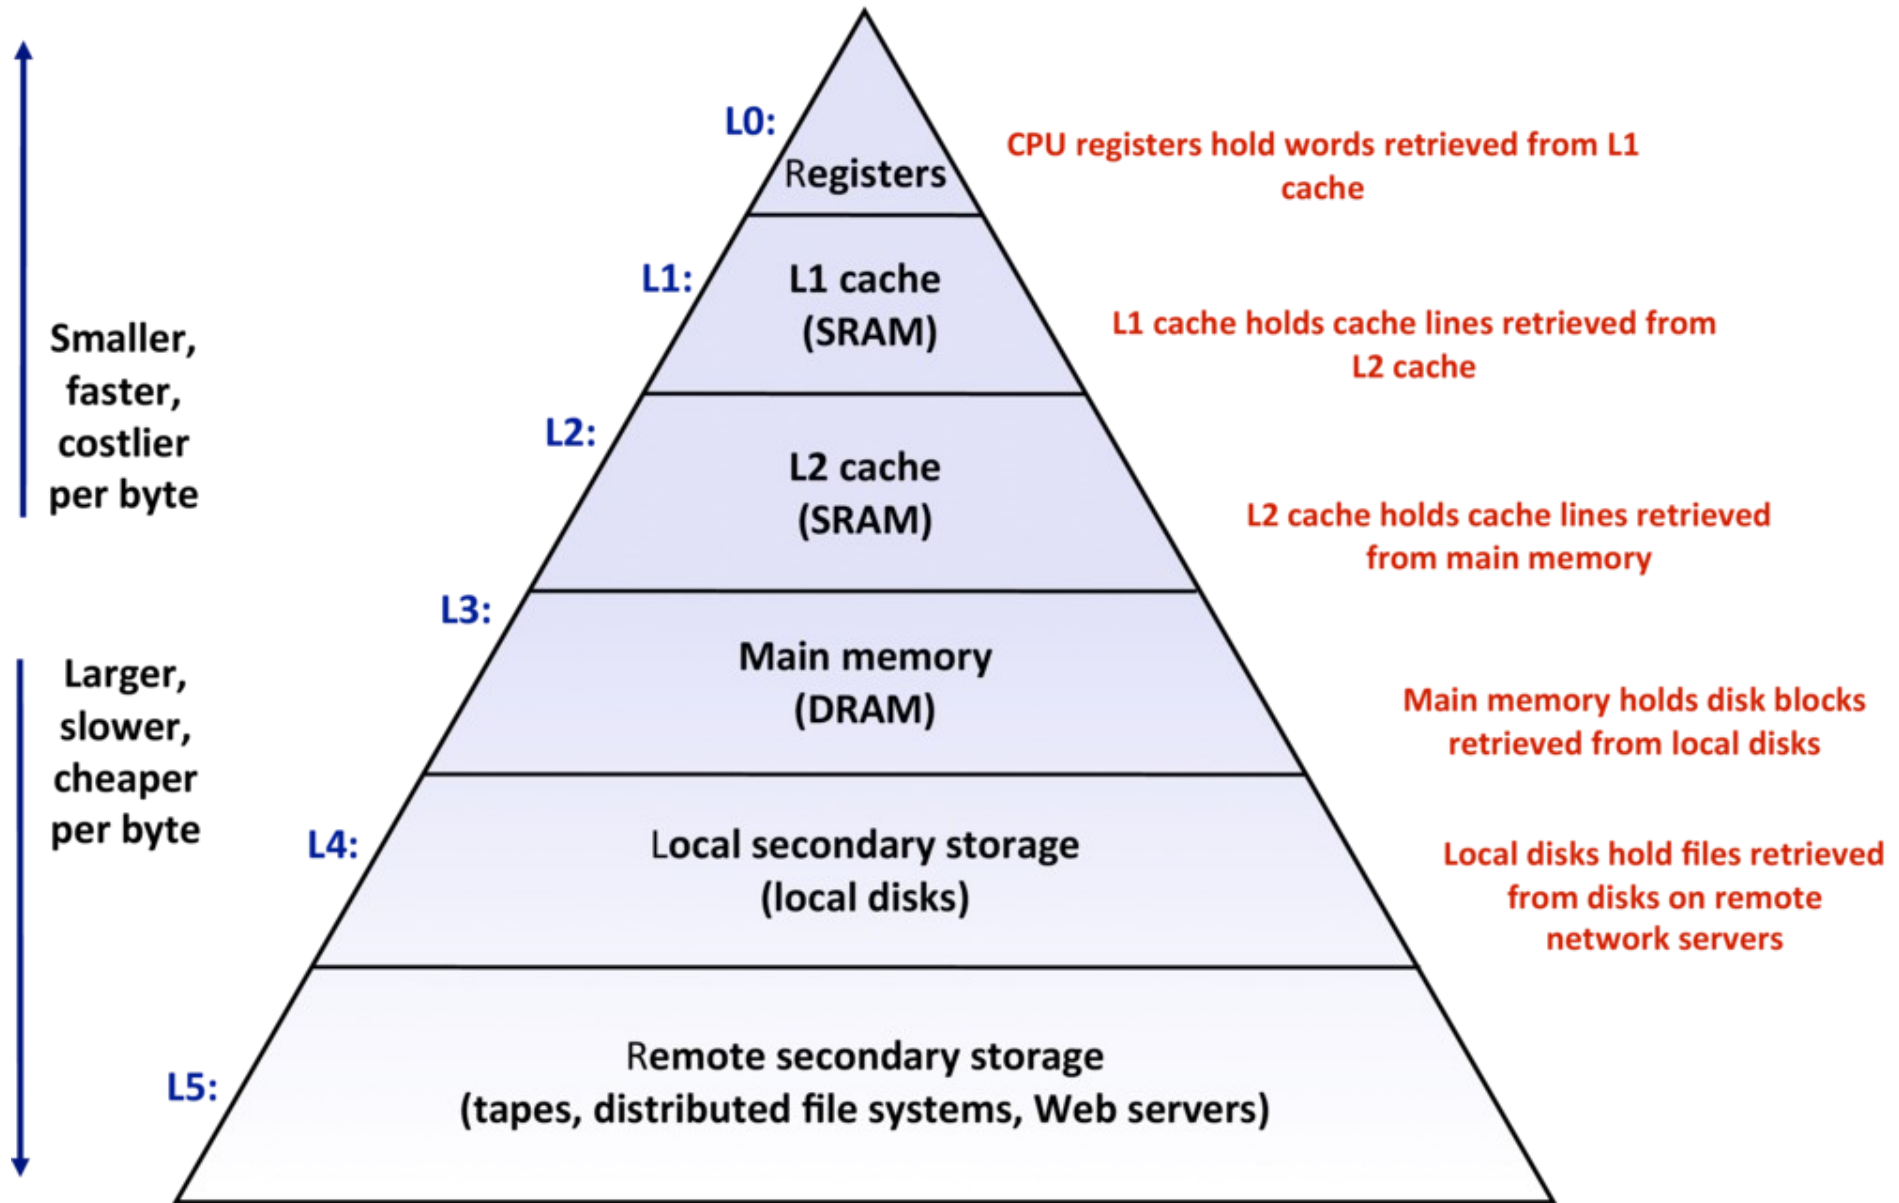


Caching

- *Cache*
 - a smaller, faster storage device that stores recently accessed data from a larger, slower device
- Fundamental concept of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$



Example memory hierarchy



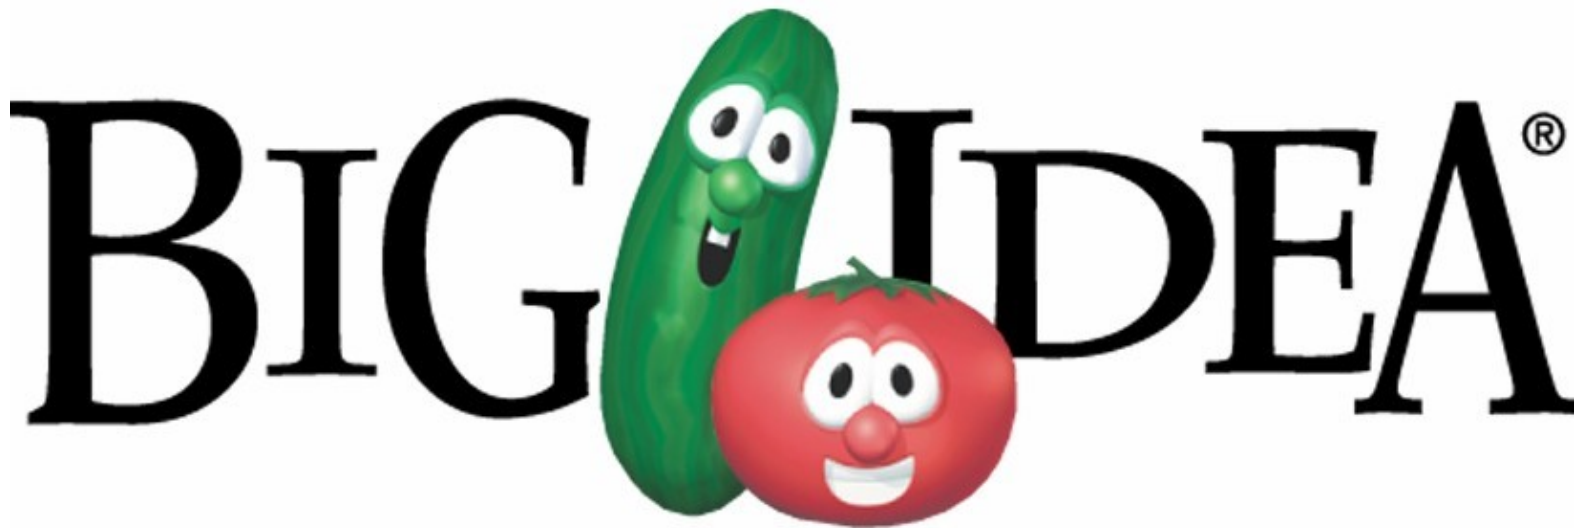
How do they work?

- Locality
 - Programs tend to access the data at level k more often than they access the data at level $k+1$.
- Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per byte.



The Big Idea

- The memory hierarchy creates a large pool of storage that **costs about as much as the cheap storage** near the bottom, but that serves data to programs **about as fast as the fast storage** near the top.



Caching example

Cache

8	9	14	3
---	---	----	---

**Smaller, faster, more expensive
memory caches a subset of
the blocks**

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
⋮			

**Larger, slower, cheaper memory
viewed as partitioned into
“blocks”**

Caching example

Cache

8	9	14	3
---	---	----	---

**Smaller, faster, more expensive
memory caches a subset of
the blocks**

**Data is copied in block-sized
transfer units**

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
⋮			

**Larger, slower, cheaper memory
viewed as partitioned into
“blocks”**

Caching example

Cache

8	9	14	3
---	---	----	---

**Smaller, faster, more expensive
memory caches a subset of
the blocks**

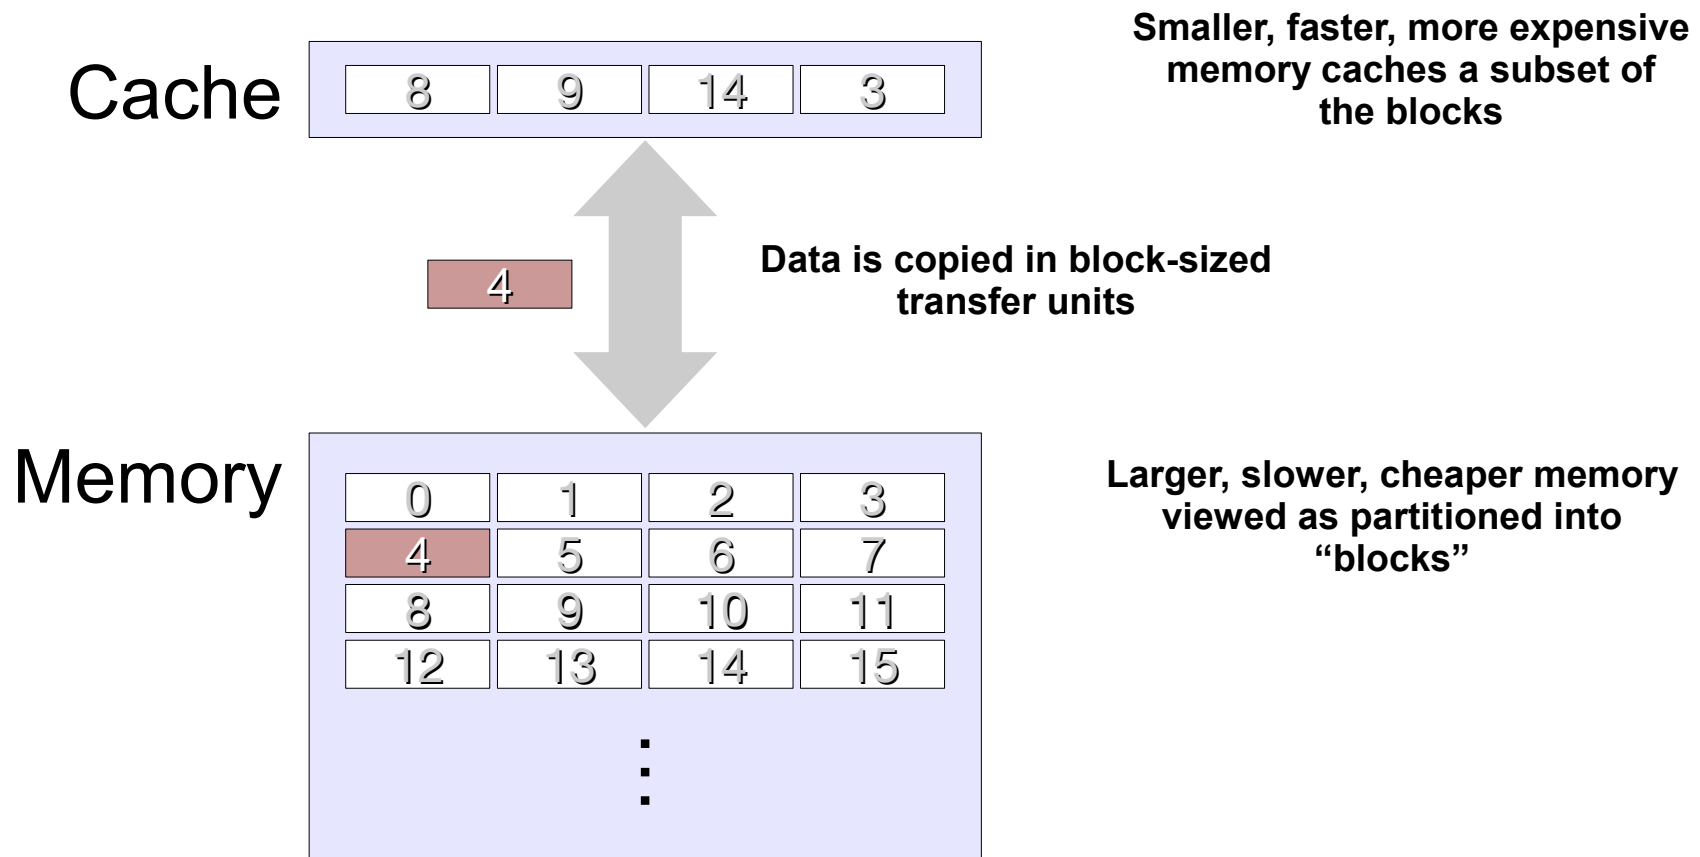
**Data is copied in block-sized
transfer units**

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
⋮			

**Larger, slower, cheaper memory
viewed as partitioned into
“blocks”**

Caching example



Caching example

Cache

4	9	14	3
---	---	----	---

**Smaller, faster, more expensive
memory caches a subset of
the blocks**

**Data is copied in block-sized
transfer units**

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
⋮			

**Larger, slower, cheaper memory
viewed as partitioned into
“blocks”**

Caching example

Cache

4	9	14	3
---	---	----	---

**Smaller, faster, more expensive
memory caches a subset of
the blocks**

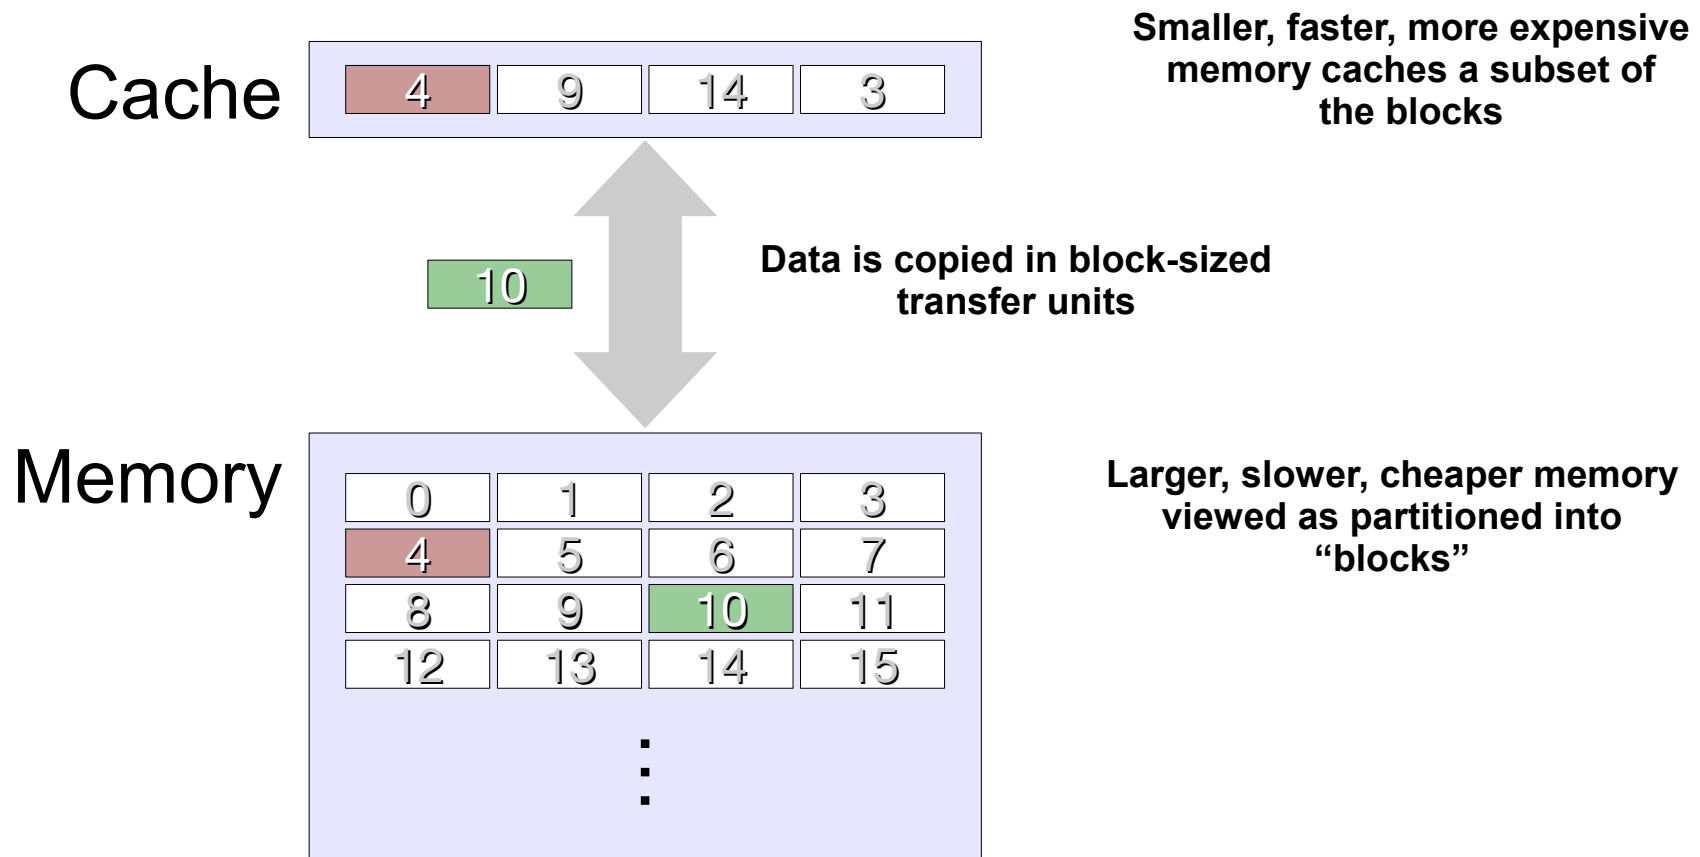
**Data is copied in block-sized
transfer units**

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
⋮			

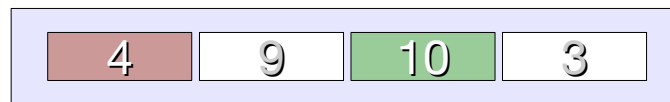
**Larger, slower, cheaper memory
viewed as partitioned into
“blocks”**

Caching example



Caching example

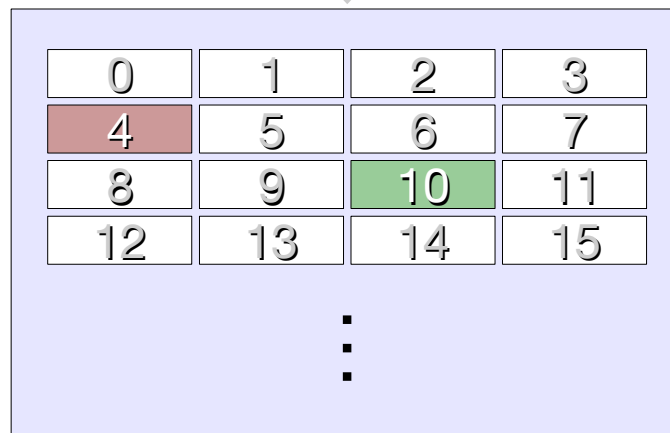
Cache



**Smaller, faster, more expensive
memory caches a subset of
the blocks**

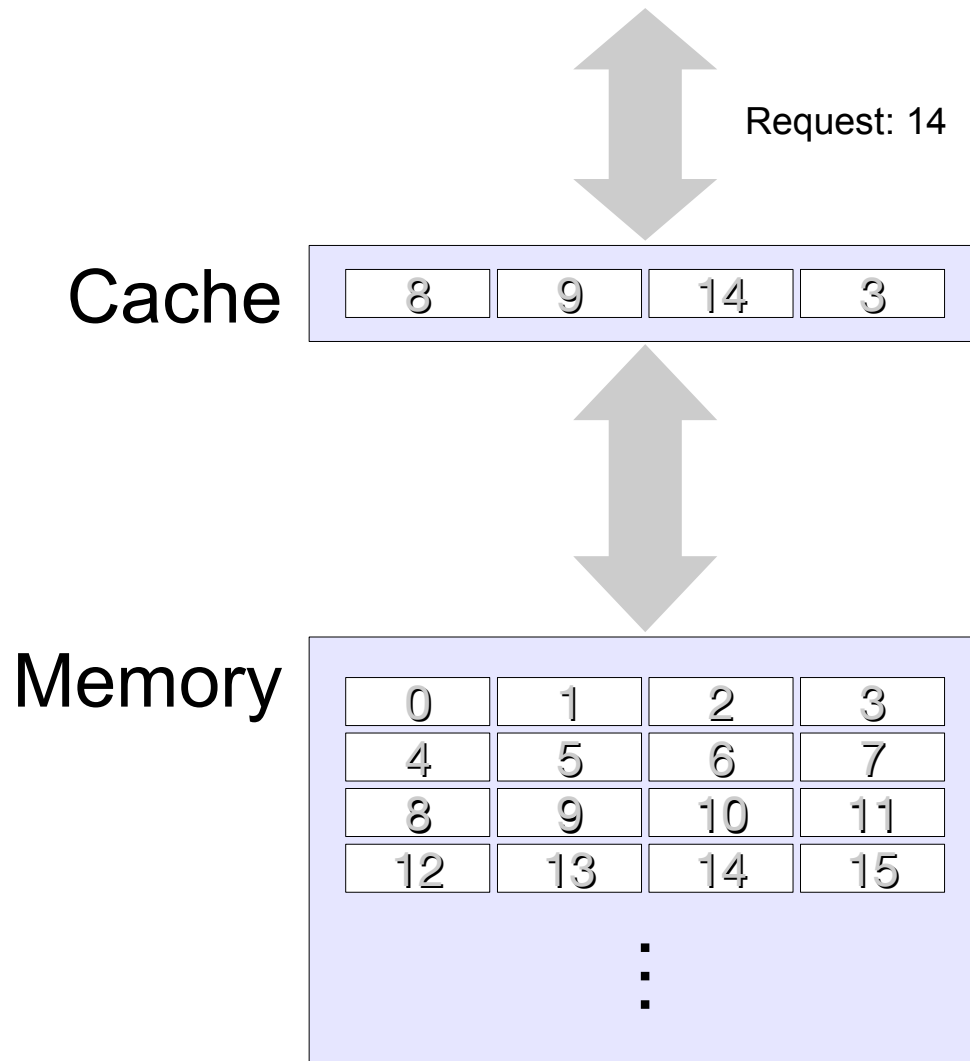
**Data is copied in block-sized
transfer units**

Memory



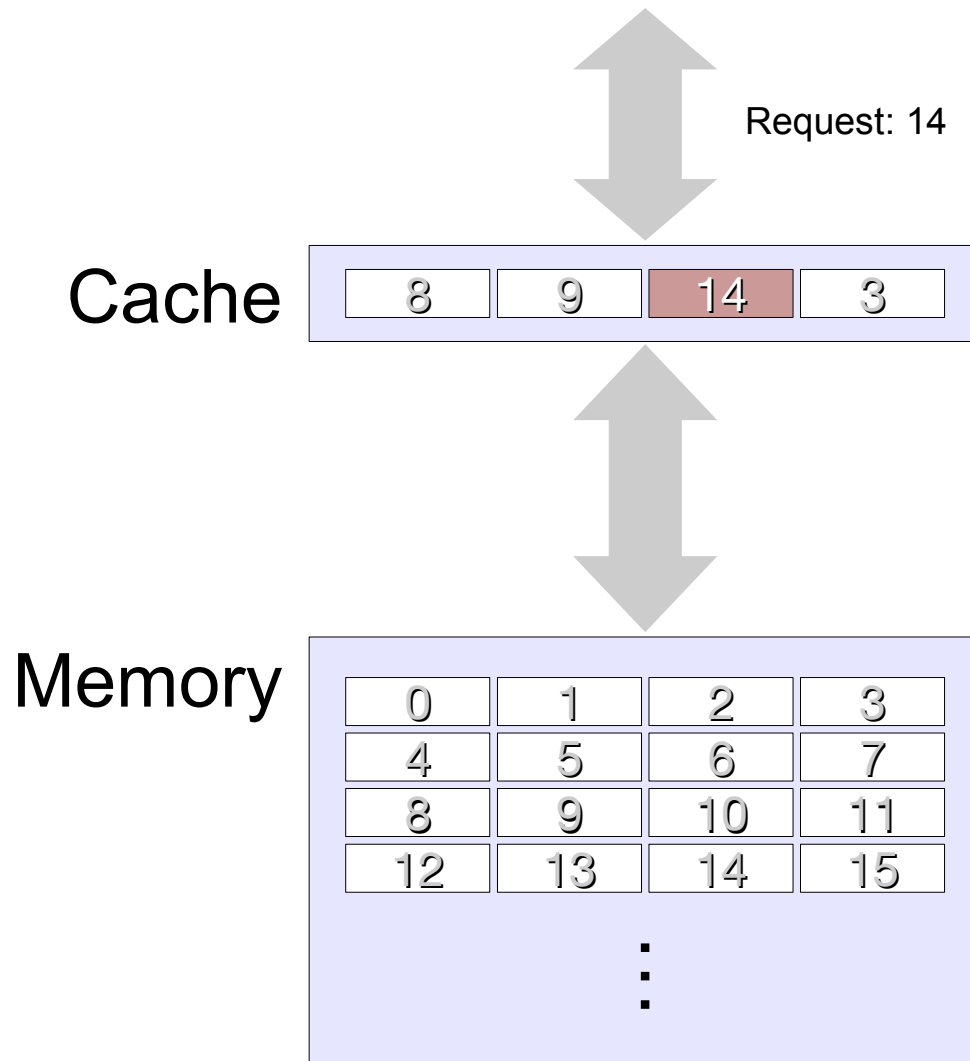
**Larger, slower, cheaper memory
viewed as partitioned into
“blocks”**

It's a hit!



Data in block *b* is needed

It's a hit!

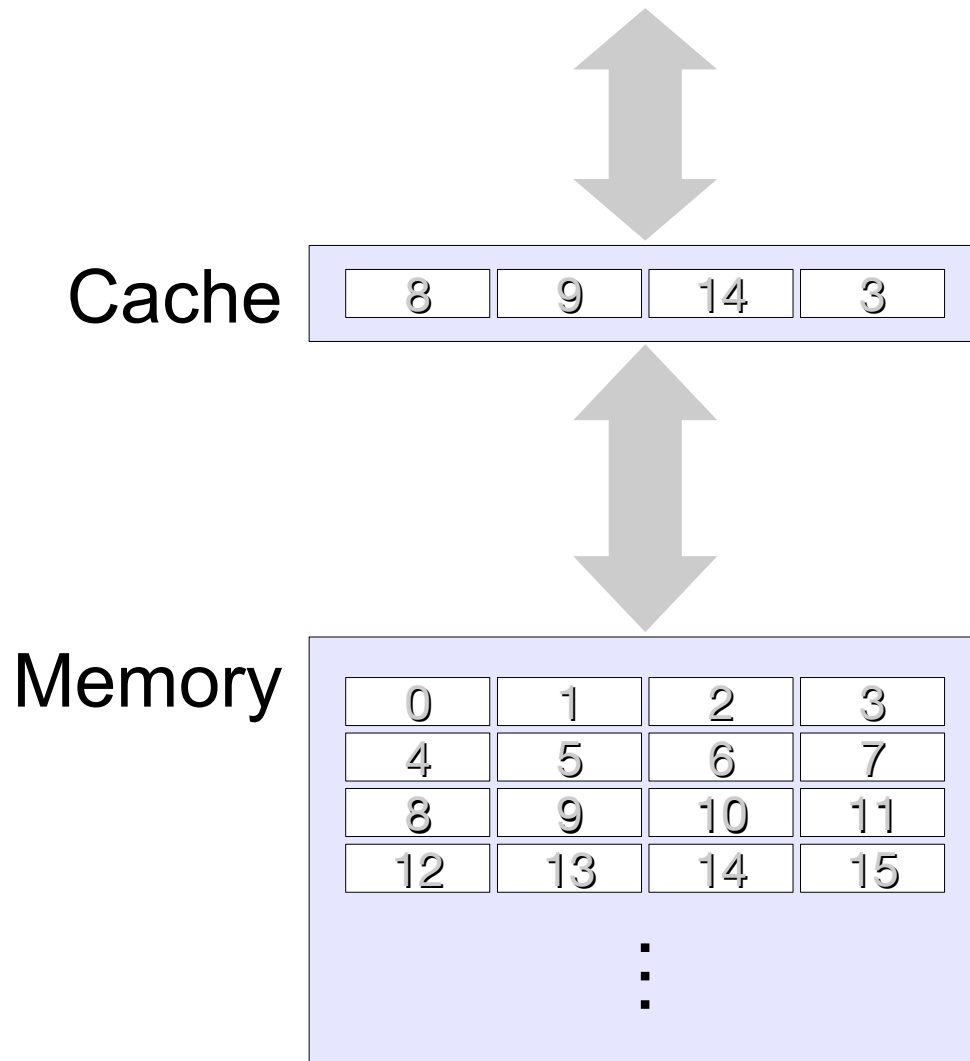


Data in block *b* is needed

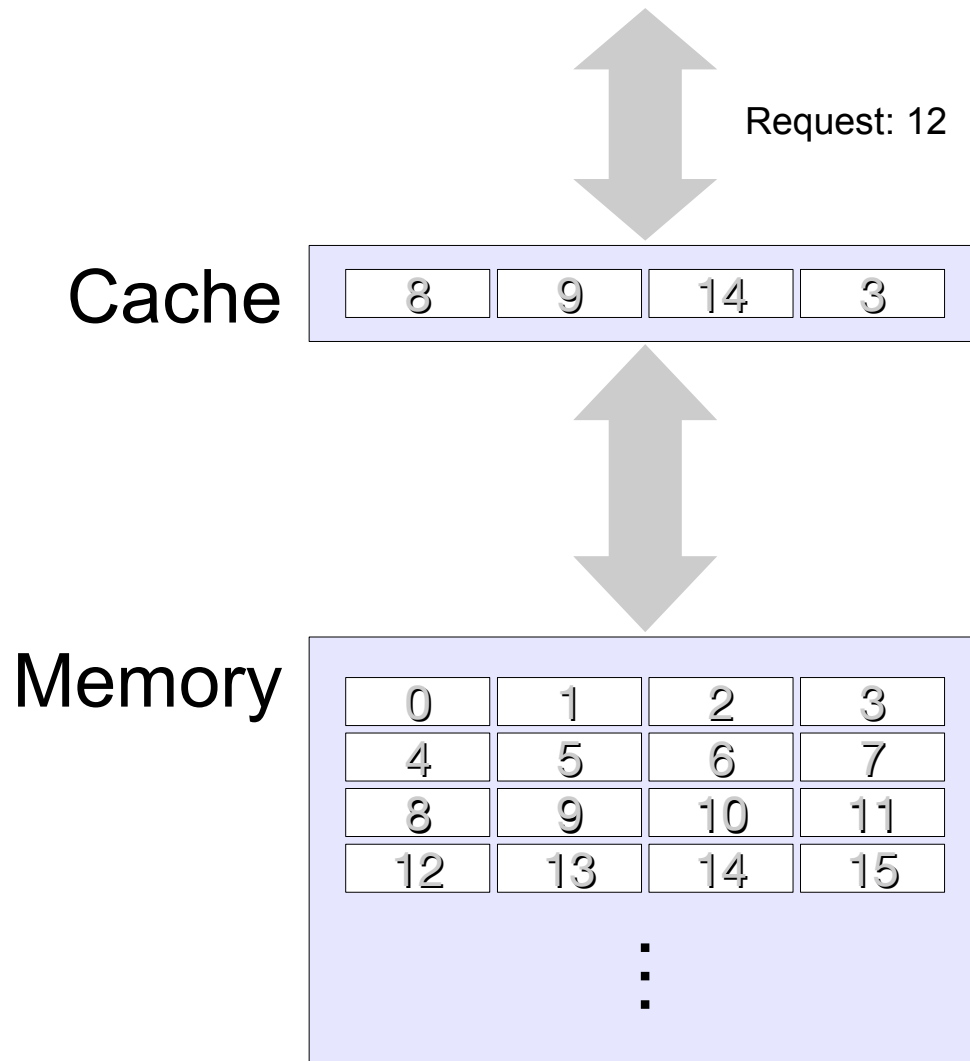
Block *b* is in cache:

Hit!

Swing and a miss

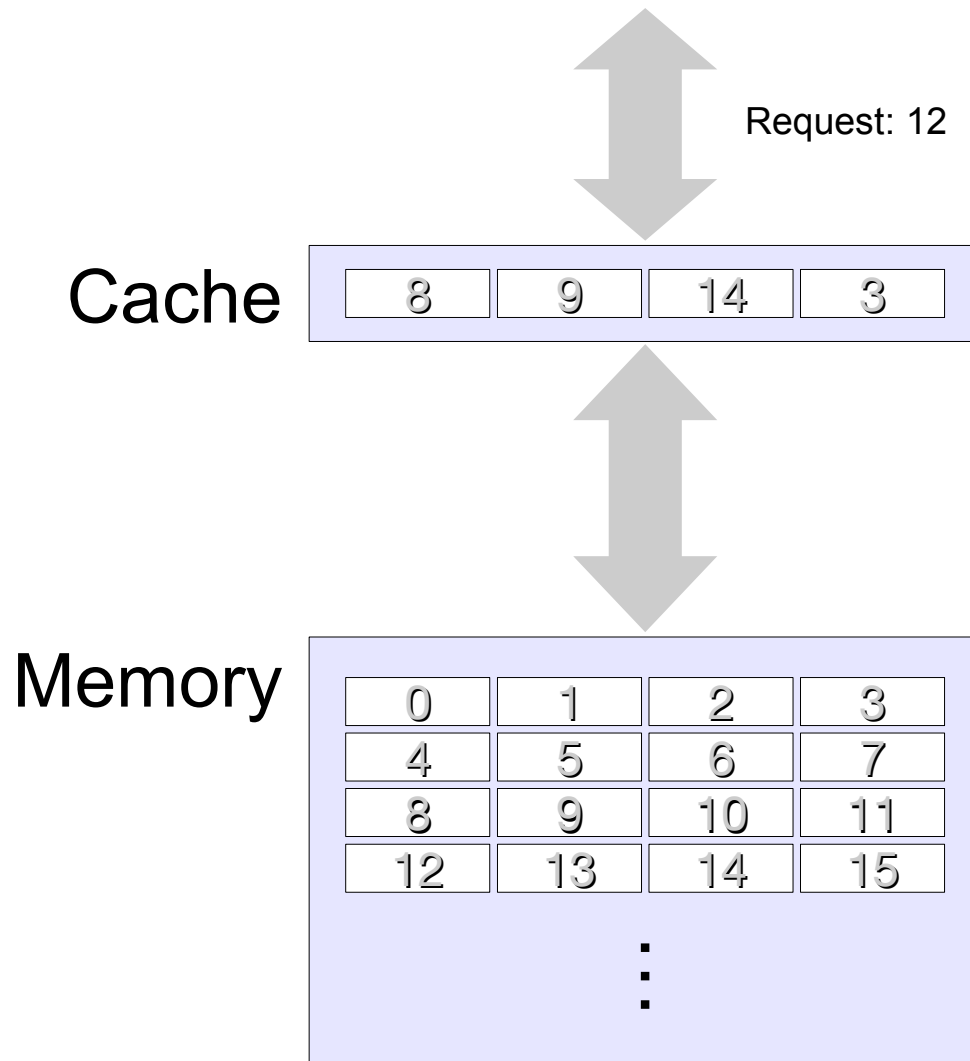


Swing and a miss



Data in block *b* is needed

Swing and a miss

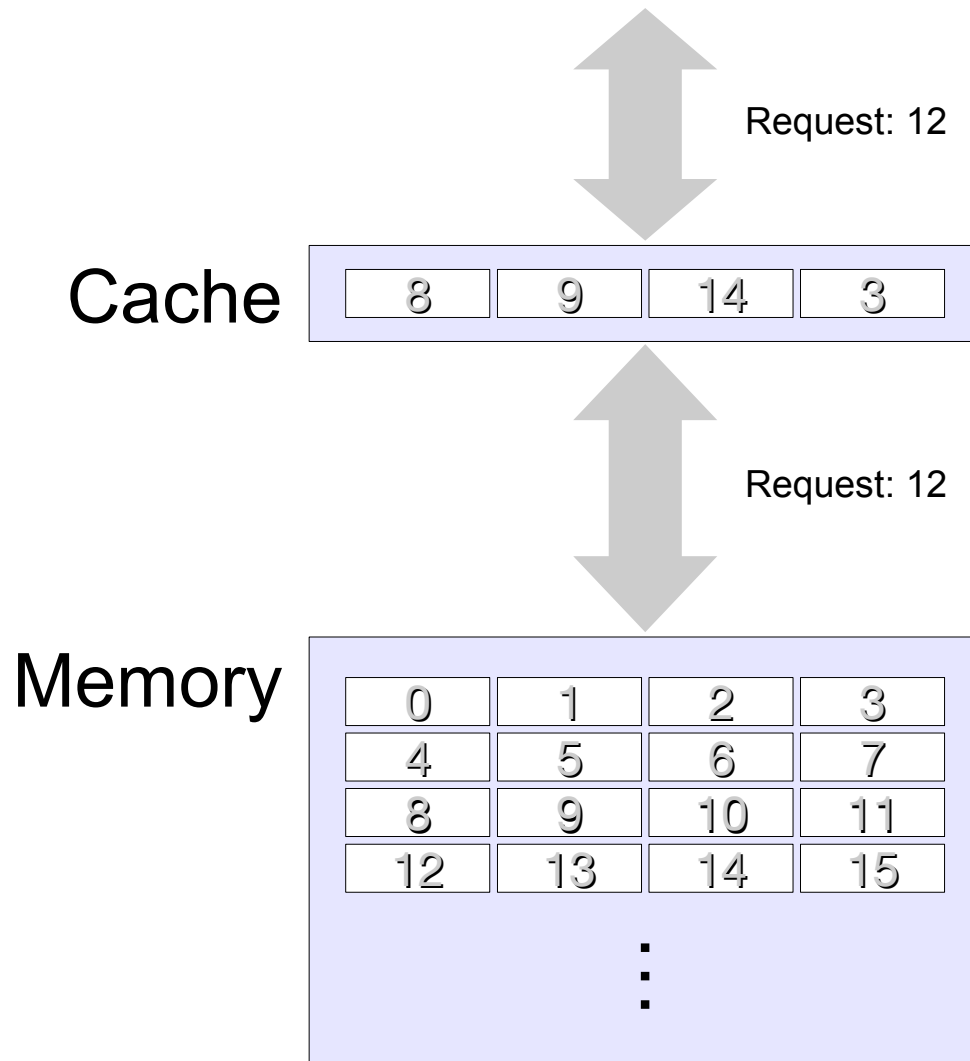


Data in block *b* is needed

Block *b* is not in cache:

Miss!

Swing and a miss



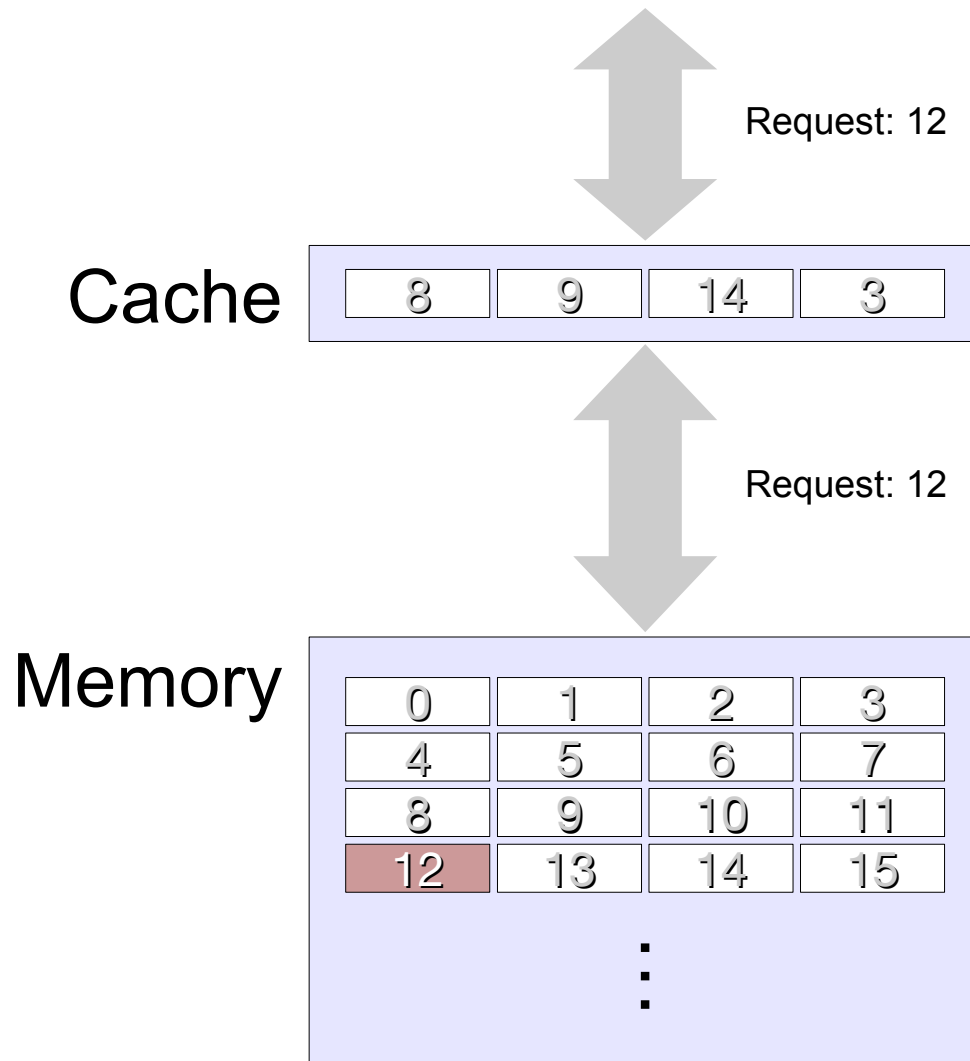
Data in block *b* is needed

Block *b* is not in cache:

Miss!

Block *b* is fetched from memory.

Swing and a miss



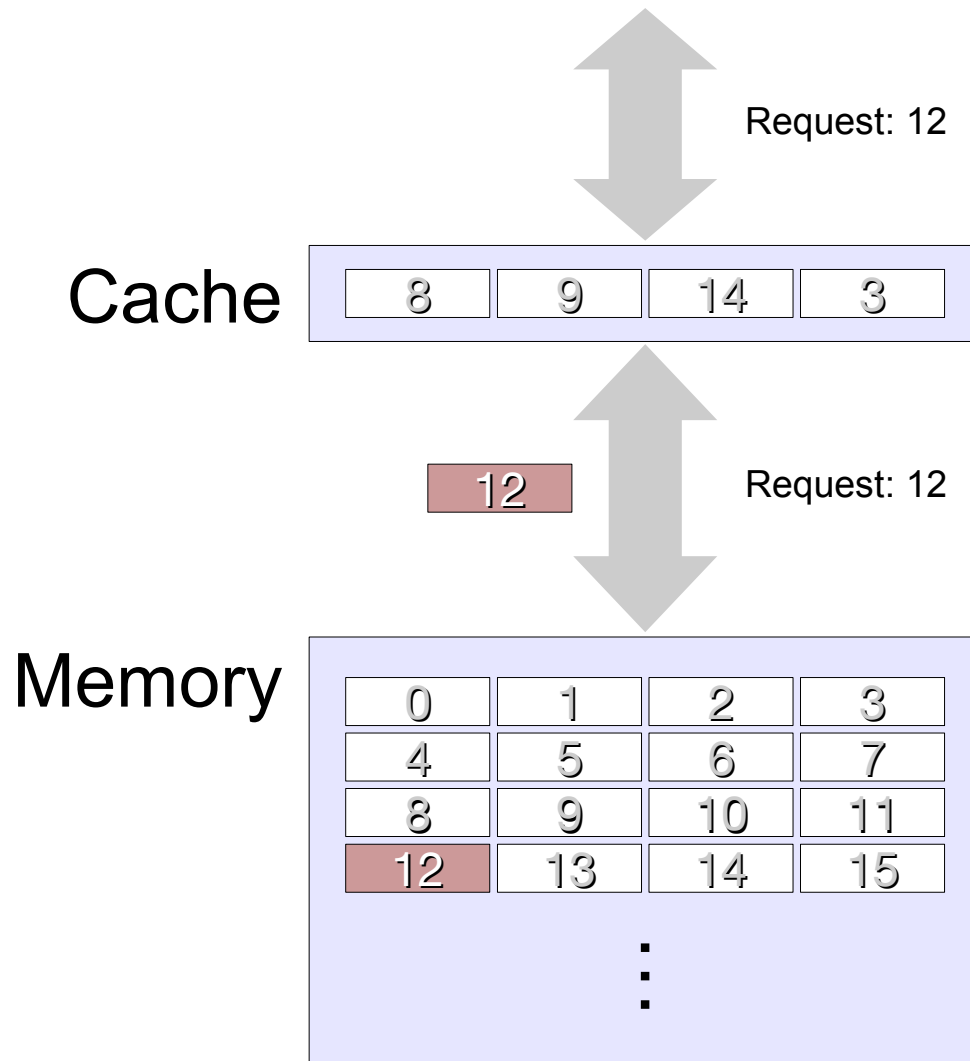
Data in block *b* is needed

Block *b* is not in cache:

Miss!

Block *b* is fetched from memory.

Swing and a miss



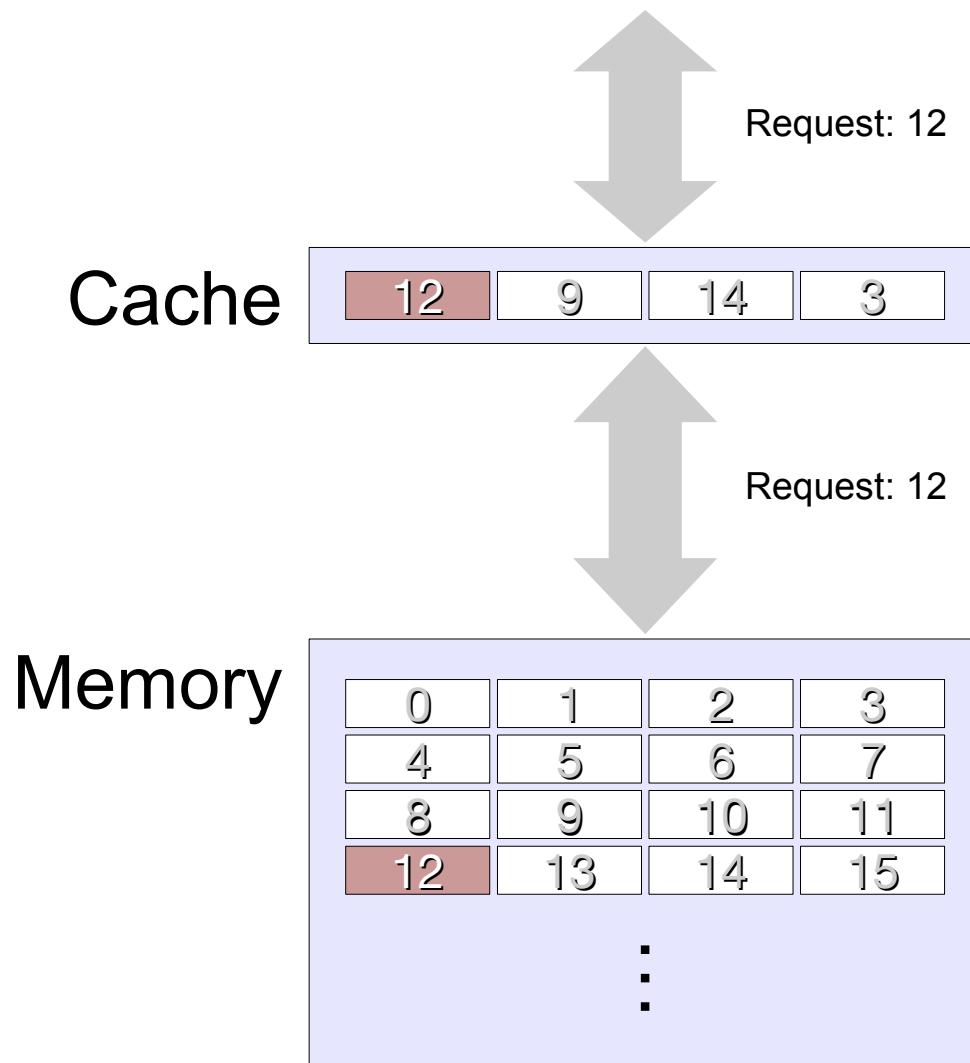
Data in block *b* is needed

Block *b* is not in cache:

Miss!

Block *b* is fetched from memory.

Swing and a miss



Data in block *b* is needed

Block *b* is not in cache:

Miss!

**Block *b* is fetched from
memory.**

Block *b* is stored in cache

Placement policy:
determines where *b* goes

Replacement policy:
determines which “victim”
block gets evicted

Reasons for cache misses

- Cold (compulsory) miss
 - Occurs because the cache is empty
- Cache is “cold” when it has just been turned on or flushed
- Cache is “hot” when it is full and likely to hit



Reasons for cache misses

- Capacity miss
 - Occurs when the cache is not big enough to hold the working set
 - **Working set**: the blocks which are currently being accessed frequently



Reasons for cache misses

- **Conflict miss**
 - Occurs when the cache is large enough, but the block has been evicted because of policy
- Most caches limit blocks from level $k+1$ to a subset (or just one) of the positions at level k .
 - Why?
 - E.g., block i in level $k+1$ must be placed at $(i \bmod 4)$ in level k .
 - In this case, referencing blocks 0, 8, 0, 8, ... would miss every time.



Cache miss example

Level k

Offset	Data
100	X
101	-
110	K

Level $k+1$

Offset (mod 8)	Data		
100	X	Y	...
101	A	B	...
110	G	K	...

Cache miss example

Level k

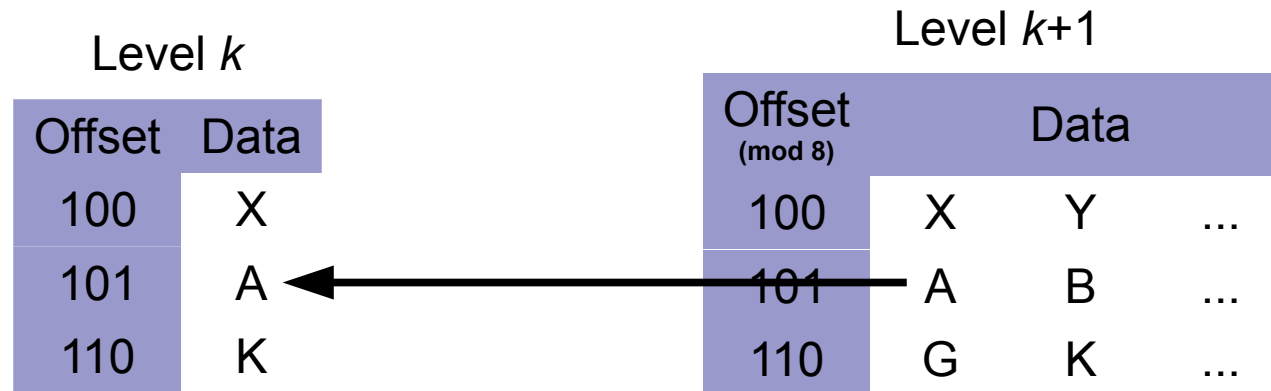
Offset	Data
100	X
101	-
110	K

Level $k+1$

Offset (mod 8)	Data		
100	X	Y	...
101	A	B	...
110	G	K	...

Request: block A

Cache miss example



Request: block A
Cold miss, fetched from level $k+1$

Cache miss example

Level k

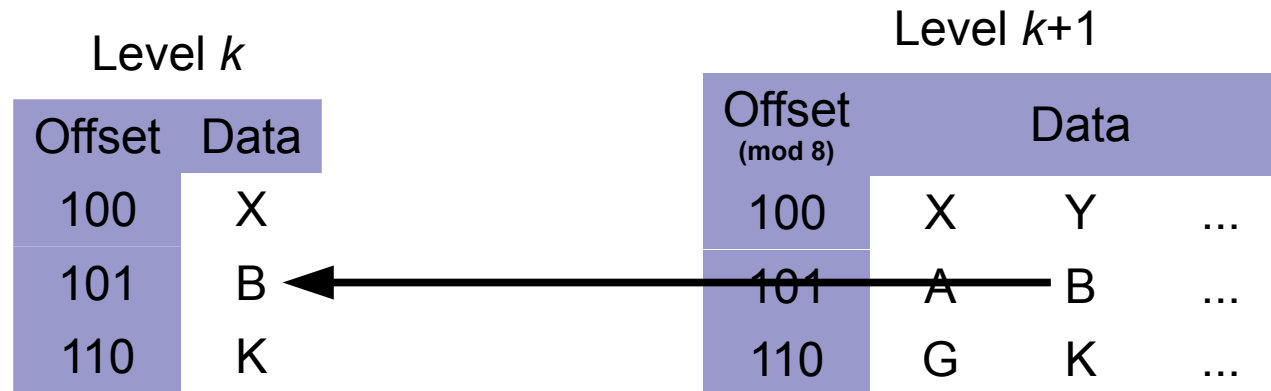
Offset	Data
100	X
101	A
110	K

Level $k+1$

Offset (mod 8)	Data		
100	X	Y	...
101	A	B	...
110	G	K	...

Request: block B

Cache miss example



Request: block B
Conflict miss, fetched from level $k+1$

Cache replacement

- When your cache is full and you acquire a new value, you must evict a previously stored value
 - Performance of cache is affected by how appropriate the cache replacement policy is
 - Popular policies
 - **Least recently used (LRU)** – evict the value that has been in the cache the longest without being accessed
 - **Least frequently used (LFU)** – evict the value that has been access the fewest number of times
 - **First in, first out (FIFO)** – evict the values in the same order they came in
 - Efficiency of a policy is measured by the hit ratio (number of cache hits vs. total memory accesses) and measured costs
 - Determined by working set and workload

Caching in the hierarchy

Cache Type	What is Cached?	Where is it Cached?	Latency (cycles)	Managed By
Registers	4-8 byte words	CPU core	~0	Compiler
TLB	Address translations	On-chip TLB	~0	Hardware
L1 cache	64-byte blocks	On-chip L1	1	Hardware
L2 cache	64-byte blocks	On/off-chip L2	10	Hardware
Virtual memory	4kiB pages	Main memory	100	Hardware + OS
Buffer cache	Parts of files	Main memory	100	OS
Disk cache	Disk sectors	Disk controller	100,000	Disk firmware
Network buffer cache	Parts of files	Local disk	10,000,000	Network FS client
Browser cache	Web pages	Local disk	10,000,000	Web browser
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

Next lecture

- Unix Essentials – getting started with using a Unix-like system
 - Basics of the command line
 - Using a traditional and powerful CLI editor
 - Writing simple scripts
- Hands-on! Bring your laptop (with VM) if you have one.



UP NEXT