

Types, Structs, and Unions

Devin J. Pohly <djpohly@cse.psu.edu>

- A *data type* is just an abstraction
 - Allows a programmer to treat a memory range as if it were a certain kind (type) of data
 - e.g., integers, real numbers, strings of characters, records, etc.
- **Note:** a variable name simply acts as an alias for a memory range.

```
// All this does is allocate memory  
// with implicit/explicit sizes and  
// values:
```


```
short int si = 9;  
long int li = 1234567890L;  
float f = 3.14;  
double d = 12324567890.1234567;  
char c = 'a';  
char *ptr = &c;
```

- The compiler uses types to determine what exactly the code does
 - How different variables can be operated on
 - Ultimately what machine instructions are generated and executed
- All programming languages use a *type system* to interpret code.

```
// Is this legal?  
double one = 3.24, two = 4.5, res1;  
int three = 3, four = 4059, res2;  
  
// Are the ISA instructions for  
// these two operations the same?  
res1 = one + two;  
res2 = three + four;
```

C typing

- Programming languages are often classified as *strongly* or *weakly typed*
 - Or somewhere in between...
 - Such distinctions refer to how the language deals with *ambiguity* in types and usage
 - C is weakly typed, which leads to great flexibility
 - ... and the potential for great bugs



- What value is output from the following code?
- Was that what the programmer intended?

```
// How is "one" treated?  
double one = 3.24;  
  
if (one / 3 == 1) {  
    printf("true");  
} else {  
    printf("false");  
}
```

Static vs. dynamic typing

- There is no clear, widely-accepted definition of strong and weak typing
- Another way to look at typing:
 - *Static typing*: the type of data is decided at compile time, and type conversion occurs at that time
 - Examples: C, C++, Java
 - *Dynamic typing*: the run-time environment dynamically applies types to variables as needed
 - Examples: Perl, Python, Ruby, Objective-C

Type coercion

- Often you want to change the type of a variable
- *Type coercion* relies on the language to do it automatically
- Coercion: the integer literals 3 and 1 are turned into double expressions, so the output is “false”. Was that the intent?
- What if you want to control the way the type is converted?

```
// How is "one" treated?  
double one = 3.24;  
  
if (one / 3 == 1) {  
    printf("true");  
} else {  
    printf("false");  
}
```

Type casting

- You can effect a temporary type conversion explicitly using *type casting*:

(type) expr

where *type* is the type to convert to, and *expr* is the variable or expression to convert

```
// How is it treated now?  
double one = 3.24;  
  
if ((int) one / 3 == 1) {  
    printf("true");  
} else {  
    printf("false");  
}
```


Legal type casting


```
int main(int argc, char *argv[]) {
    short int si = 9;
    long int li = 1234567890L;
    float f = 3.14;
    double d = 12324560.1234567;
    char c = 'a';
    char *ptr = &si;

    printf("short int %d %f %p\n", (int) si, (float) si, (char *) si);
    printf("long int %d %f %p\n", (int) li, (float) li, (char *) li);
    printf("float %d %f (ERR)\n", (int) f, (float) f);
    printf("double %d %f (ERR)\n", (int) d, (float) d);
    printf("char %d %f %p\n", (int) c, (float) c, (char *) c);
    printf("ptr %d (ERR) %p\n", (int) ptr, (char *) ptr);
    return 0;
}
```

```
short int 9 9.000000 0x9
long int 1234567890 1234567936.000000 0x499602d2
float 3 3.140000 (ERR)
double 12324560 12324560.000000 (ERR)
char 97 97.000000 0x61
ptr -716365630 (ERR) 0x7fffd54d20c2
```

Structure types

- A struct is an organized set of data that are treated as a unit
 - I.e., it behaves like a single variable
 - Can be declared like any other variable
 - Type name includes the keyword struct and a tag you provide



Struct syntax and use

```
// Define the struct
struct myname {
    type1 field1;
    type2 field2;
    ...
};

// Declare/allocate a struct variable
struct myname foo;

// Access a field
x = foo.field1;
foo.field1 = 42;
```

- Here, typeX is the type of each piece of data (*field*) inside this struct, and fieldX is the name of that field
 - Note the semicolon after the braces! This is a *statement* that defines a new type.


Basic example

```
struct vehicle {  
    // Make and model  
    char name[128];  
  
    // Current mileage  
    int mileage;  
};  
  
struct vehicle  
    cayman, gremlin,  
    cessna180, montauk;  
  
...  
  
montauk.mileage = 5500;  
  
printf("%s: %d mi.\n",  
    gremlin.name,  
    gremlin.mileage);
```



Enumerations

- Allows you to treat a **set of names** as an efficient integer value
 - Used when there is a discrete range of values for a type that **isn't already numeric**
 - Type name includes the keyword `enum` and a tag



Enumeration syntax

```
enum vehicle_type {  
    AUTOMOTIVE,  
    AERONAUTICAL,  
    MARINE  
};  
  
enum vehicle_type t;  
  
t = MARINE;
```

- Note again the semicolon!
- You can also assign specific values if needed, e.g.:
 SPECIAL = 999

Nested types

- Structs can include fields of any known type
 - Primitive types like int or float
 - Arrays
 - Pointers
 - **Other structs!**
 - Enumerations
 - etc.

```
struct engine_specs {  
    int cylinders;  
    int horsepower;  
};  
  
enum vehicle_type {  
    AUTOMOTIVE,  
    AERONAUTICAL,  
    MARINE  
};  
  
struct vehicle {  
    // Make and model  
    char name[128];  
    // Current mileage  
    int mileage;  
    // Type of vehicle  
    enum vehicle_type type;  
    // Engine specifications  
    struct engine_specs engine;  
};
```

- A union allows you to use the *same memory region* for several variables
 - Can even be of different types – **DANGER!**
 - C will not stop you from writing bits as one type and reading them as a different one!
 - If you ever use a union, you should consider its members to be *mutually exclusive!*




```
// Define the union
union vehicle_ident {
    char vin[17];
    char tail_number[8];
    char hull_id[12];
};

// Declare/allocate a union variable
union vehicle_ident id;

// Access a member
printf("VIN: %s\n", id.vin);
```

- Defined just like a struct
 - Instead of containing **all of the fields**, it contains **one of the fields**
 - What type it treats the memory as depends on which member you are accessing
 - Only needs to be as big as the largest member

Bringing it all together

```
struct engine_specs {
    int cylinders;
    int horsepower;
};

enum vehicle_type {
    AUTOMOTIVE,
    AERONAUTICAL,
    MARINE
};

union vehicle_ident {
    char vin[17];
    char tail_number[8];
    char hull_id[12];
};

struct vehicle {
    char name[128];
    int mileage;
    enum vehicle_type type;
    struct engine_specs engine;
    union vehicle_ident id;
};

struct vehicle cayman, gremlin, cessna180, montauk;
```

Type aliases: typedef

- Type names sometimes get long and unwieldy
- The C typedef statement creates an *alias* for an existing type:

```
typedef oldtype newtype;
```

- where
 - *oldtype* is a type name, written exactly as if you were declaring a variable
 - *newtype* is the new alias for this type
 - Convention: type aliases frequently end in `_t`
- Example:

```
typedef struct engine_specs engine_t;
```

Using user-defined types

- You can use the new type anywhere you use built-in types
 - Provides an abstraction
- Note: the compiler treats the alias exactly as if it were the original type

```
// Type declaration
typedef unsigned char age_t;

// Return values and function parameters
age_t myFunction(age_t x, int y) {
    // Local variables
    age_t z;
    float a;


    // Type casting
    return (age_t) 1;
}
```

Struct pointers

- Work exactly like any other pointer

```
struct vehicle *pv = &cayman;
```

- How could you get the mileage field given this pointer?




Struct pointers

- Work exactly like any other pointer

```
struct vehicle *pv = &cayman;
```

- How could you get the mileage field given this pointer?

```
int m = (*pv).mileage;
```



Struct pointers

- Work exactly like any other pointer


```
struct vehicle *pv = &cayman;
```

- How could you get the mileage field given this pointer?

```
int m = (*pv).mileage;
```

- So common that it has its own syntactic sugar: the arrow operator.

```
int m = pv->mileage;
```



Struct layout in memory

```
#define MEM_OFFSET(a, b) (((uintptr_t) &b) - ((uintptr_t) &a))

// Print out the addresses of the fields
printf("          SZ      Addr          ofs\n");
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman), &cayman, MEM_OFFSET(cayman, cayman));
printf("cayman.name     %3lu %p 0x%02lx\n",
      sizeof(cayman.name), &cayman.name, MEM_OFFSET(cayman, cayman.name));
printf("cayman.mileage    %3lu %p 0x%02lx\n",
      sizeof(cayman.mileage), &cayman.mileage, MEM_OFFSET(cayman, cayman.mileage));
printf("cayman.type       %3lu %p 0x%02lx\n",
      sizeof(cayman.type), &cayman.type, MEM_OFFSET(cayman, cayman.type));
printf("cayman.engine.cylinders %3lu %p 0x%02lx\n",
      sizeof(cayman.engine.cylinders), &cayman.engine.cylinders,
      MEM_OFFSET(cayman, cayman.engine.cylinders));
printf("cayman.engine.horsepower %3lu %p 0x%02lx\n",
      sizeof(cayman.engine.horsepower), &cayman.engine.horsepower,
      MEM_OFFSET(cayman, cayman.engine.horsepower));
printf("cayman.vehicle_id.vin %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.vin), &cayman.vehicle_id.vin,
      MEM_OFFSET(cayman, cayman.vehicle_id.vin));
printf("cayman.vehicle_id.tail_number %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.tail_number), &cayman.vehicle_id.tail_number,
      MEM_OFFSET(cayman, cayman.vehicle_id.tail_number));
printf("cayman.vehicle_id.hull_id %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.hull_id), &cayman.vehicle_id.hull_id,
      MEM_OFFSET(cayman, cayman.vehicle_id.hull_id));
```


Struct layout in memory

```
#define MEM_OFFSET(a, b) (((uintptr_t) &b) - ((uintptr_t) &a))

// Print out the addresses of the fields
printf("          SZ      Addr      Ofc\n");
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman), &cayman, MEM_OFFSET(cayman, cayman));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.name), &cayman.name, MEM_OFFSET(cayman, cayman.name));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.mileage), &cayman.mileage, MEM_OFFSET(cayman, cayman.mileage));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.type), &cayman.type, MEM_OFFSET(cayman, cayman.type));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.engine.cylinders), &cayman.engine.cylinders, MEM_OFFSET(cayman, cayman.engine.cylinders));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.engine.horsepower), &cayman.engine.horsepower, MEM_OFFSET(cayman, cayman.engine.horsepower));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.vin), &cayman.vehicle_id.vin, MEM_OFFSET(cayman, cayman.vehicle_id.vin));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.tail_number), &cayman.vehicle_id.tail_number, MEM_OFFSET(cayman, cayman.vehicle_id.tail_number));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.hull_id), &cayman.vehicle_id.hull_id, MEM_OFFSET(cayman, cayman.vehicle_id.hull_id));
printf("cayman          %3lu %p 0x%02lx\n",
      sizeof(cayman.vehicle_id.hull_id), &cayman.vehicle_id.hull_id, MEM_OFFSET(cayman, cayman.vehicle_id.hull_id));
```

	SZ	Addr	Ofc
cayman	160	0x601080	0x00
cayman.name	128	0x601080	0x00
cayman.mileage	4	0x601100	0x80
cayman.type	4	0x601104	0x84
cayman.engine.cylinders	1	0x601108	0x88
cayman.engine.horsepower	2	0x60110a	0x8a
cayman.vehicle_id.vin	17	0x60110c	0x8c
cayman.vehicle_id.tail_number	8	0x60110c	0x8c
cayman.vehicle_id.hull_id	12	0x60110c	0x8c

Wait a second...

- Add up the sizes of the variables:
 - $128 + 4 + 4 + 1 + 2 + \max(17, 8, 12) = 156$
 - But `sizeof(cayman)` said 160!
 - What happened?

	SZ	Addr	Ofs
cayman	160	0x601080	0x00
cayman.name	128	0x601080	0x00
cayman.mileage	4	0x601100	0x80
cayman.type	4	0x601104	0x84
cayman.engine.cylinders	1	0x601108	0x88
cayman.engine.horsepower	2	0x60110a	0x8a
cayman.vehicle_id.vin	17	0x60110c	0x8c
cayman.vehicle_id.tail_number	8	0x60110c	0x8c
cayman.vehicle_id.hull_id	12	0x60110c	0x8c

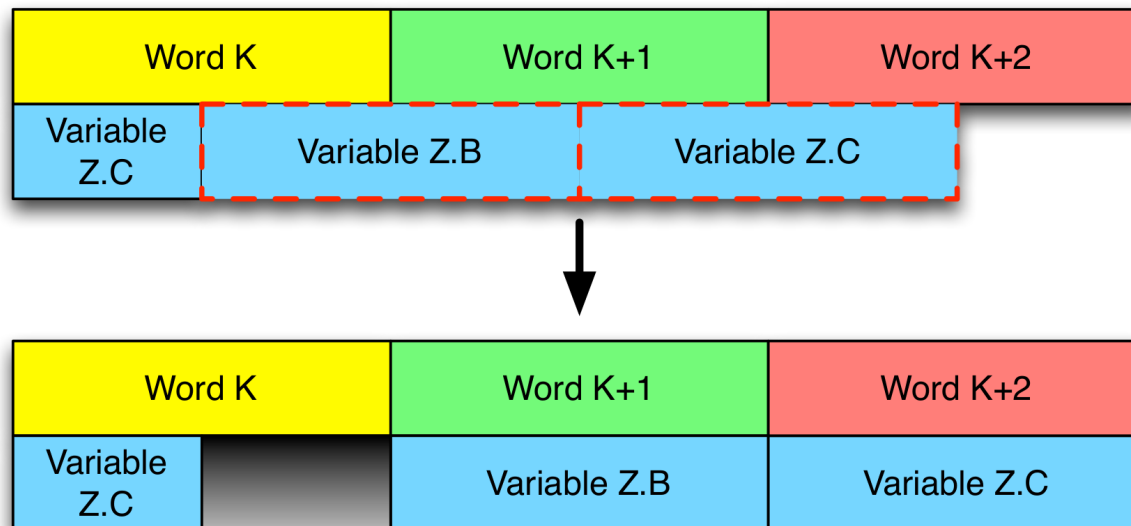
Let's take a closer look

- If we do a little hex math...
 - From offset 0x88 to 0x8a is 2 bytes (skipped 1)
 - From offset 0x8c to 160 is 20 bytes (skipped 3)

	SZ	Addr	Ofs
cayman	160	0x601080	0x00
cayman.name	128	0x601080	0x00
cayman.mileage	4	0x601100	0x80
cayman.type	4	0x601104	0x84
cayman.engine.cylinders	1	0x601108	0x88
cayman.engine.horsepower	2	0x60110a	0x8a
cayman.vehicle_id.vin	17	0x60110c	0x8c
cayman.vehicle_id.tail_number	8	0x60110c	0x8c
cayman.vehicle_id.hull_id	12	0x60110c	0x8c

The answer

- The compiler may “**pad**” the structure with unused memory so that offsets align with a multiple of the machine word size.
 - This is because many ISAs have instructions that require the target address(es) to be “**word-aligned**”
 - **Caution**: the way this works varies between architectures and compilers. Beware when working with data from other computers (network communication, file formats, etc.)!



- Often you want to create numeric (integer) fields that have a very specific width (in bits)
 - C supports this by explicitly identifying the bit width in declarations of integer fields:

```
// Define a structure of bit fields
struct vehicle_props {
    uint16_t registered: 1;
    uint16_t color_code: 8;
    uint16_t doors: 3;
    uint16_t year: 16;
};

struct vehicle_props props;

// Use the fields
props.registered = 1;
props.color_code = 14;
props.doors = 2;
props.doors = 9; // Legal, but out of range!
props.year = 2014;
```

eof(1)

