

Systems and Internet Infrastructure Security

Institute for Networking and Security Research
Department of Computer Science and Engineering
Pennsylvania State University, University Park, PA



Memory Management (Part 2)

Devin J. Pohly <djpohly@cse.psu.edu>

Scope

- The *scope* of a variable is the portion of a program where it can be accessed by name
- C has *lexical scope*
 - This means scope is known at compile-time
 - In C, the scope of a variable is always the innermost {} block (the whole file for globals)

```
int i;  
void f(int i)  
{  
    i = 1;  
}  
void g(void)  
{  
    int i = 2;  
    if (i > 0) {  
        int i;  
        i = 3;  
    }  
    i = 4;  
}  
void h(void)  
{  
    int i = 5;  
}
```

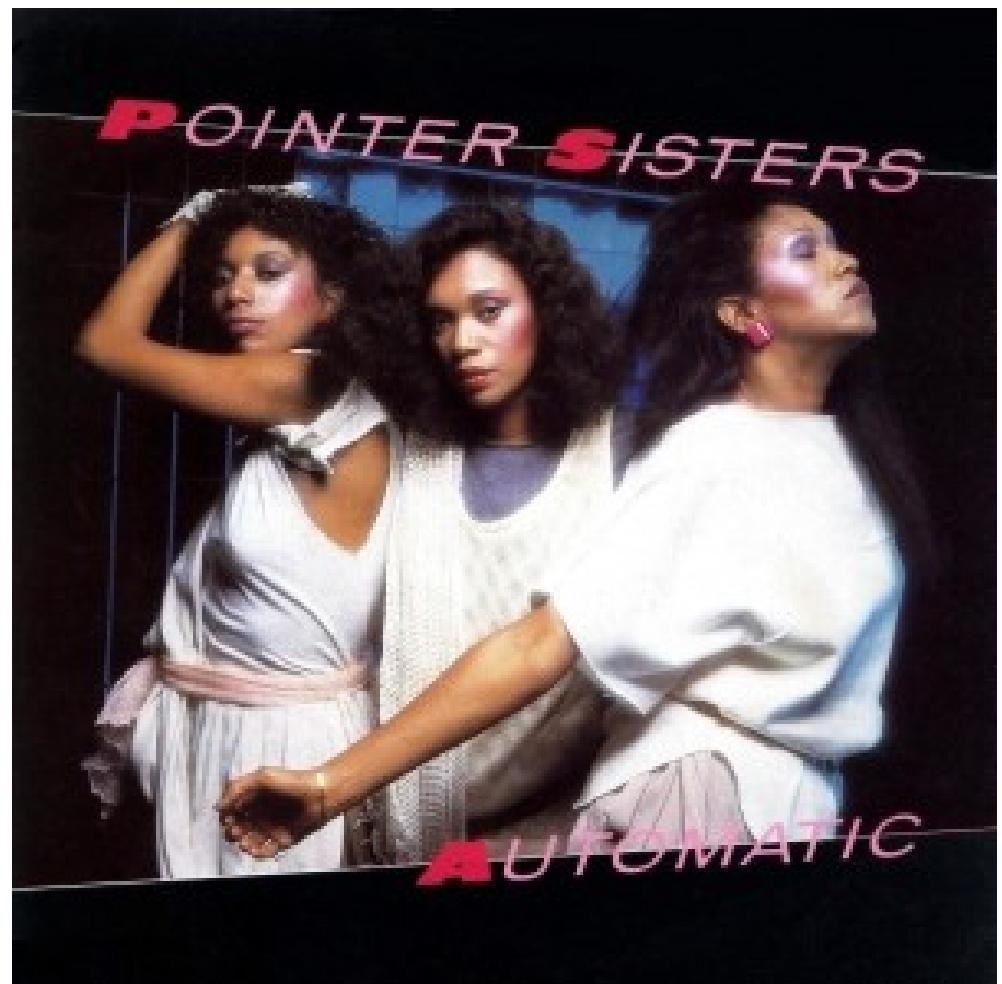
Scope vs. lifetime

- The *lifetime* of a variable is how long the data is valid (allocated)
- Different from scope
 - Scope is about the *name* of a variable
 - Lifetime is about the *data* allocated for it



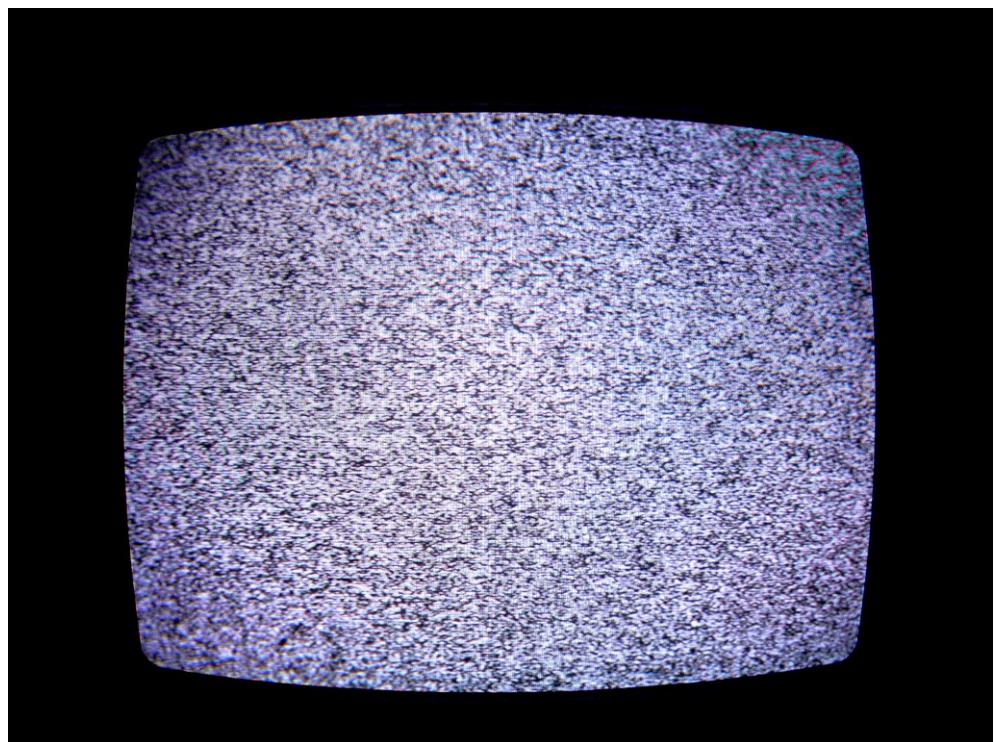
Automatic allocation

- Used for arguments and “normal” local variables
 - Generally allocated on the stack
- Scope: current function or block
- Lifetime: until the end of the function or block
 - Allocated when a function is called
 - Deallocated when it returns



Static allocation

- Used for global variables or anything else marked “static”
 - Generally allocated in the data segment
- Scope: current function or block
- Lifetime: entire life of the program
 - Allocated when the program is loaded
 - Deallocated when it exits



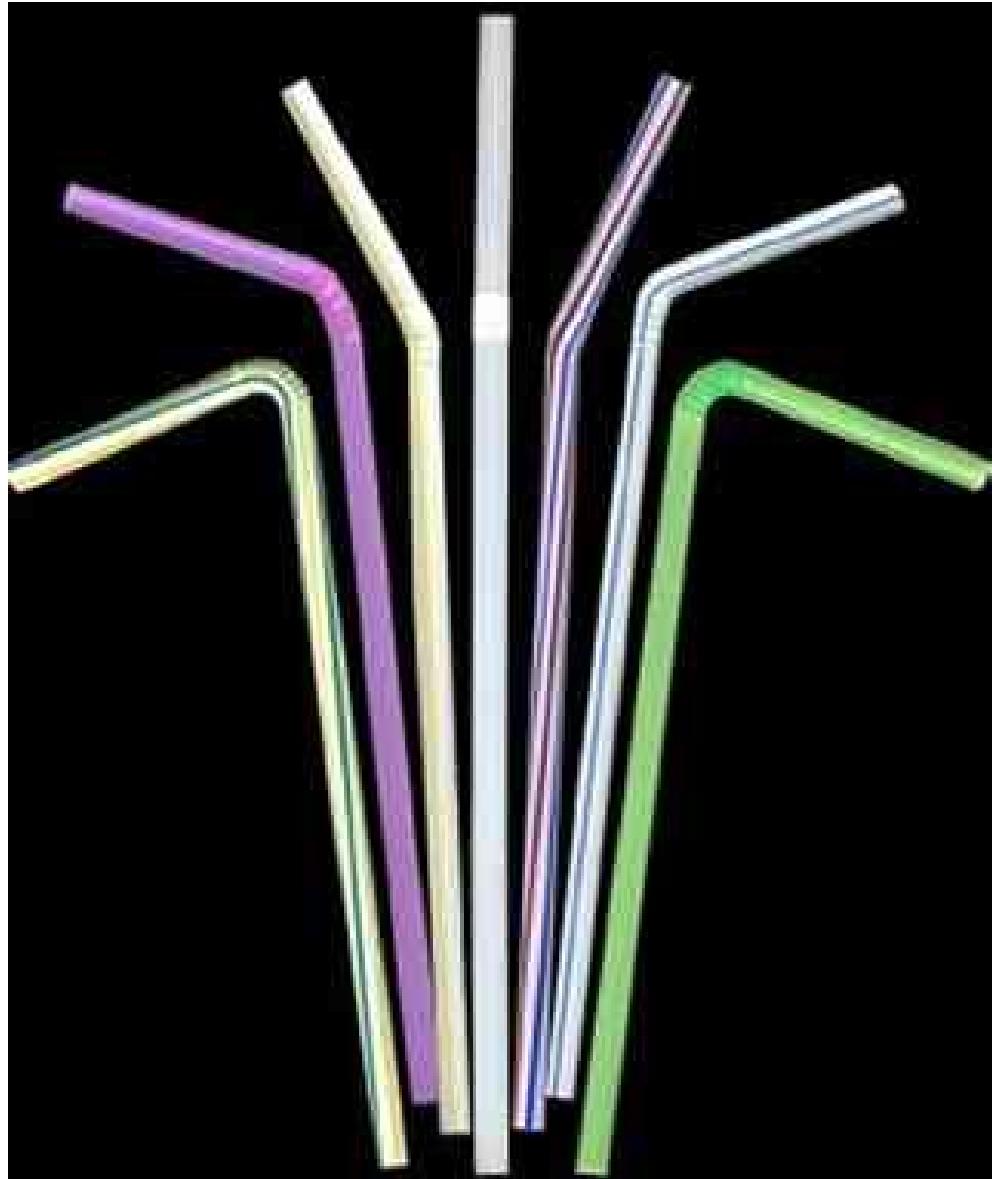
Automatic vs. static

- Which variables here are automatically allocated?
- Which are statically allocated?
- What are the scope and lifetime of the **total** variable?

```
int counter = 0;  
  
double running_avg(int value) {  
    static int total = 0;  
    double avg;  
  
    total += value;  
    counter++;  
  
    avg = (double) total / counter;  
    return avg;  
}
```

More flexibility

- What if we need memory that is...
 - valid for multiple function calls, but not the entire program?
 - too big to allocate on the stack?
 - variable-size, or the size isn't known at compile time?
 - allocated and returned by a function?



Dynamic allocation

- What we want is *dynamically allocated memory*
 - Your program explicitly requests a new block of memory
 - The C library allocates it, perhaps with help from the OS
 - Where in memory can this be allocated?
 - Dynamically allocated memory is valid until:
 - Your code explicitly deallocates it (*manual memory management*)
 - A garbage collector collects it (*automatic memory management*)
 - C requires you to manually manage memory
 - Gives you more control, but more opportunities to slip up... as usual
- Note: this essentially allows you to create custom scope and lifetime

The heap

- Large pool of memory that is used for dynamic allocation
 - Both allocated data and allocation metadata are stored here
 - Use `malloc` to allocate heap data and `free` to deallocate
 - Allocation functions are found in `<stdlib.h>`
- (Note: Unrelated to the “heap” data structure)



Allocating memory: malloc

- malloc allocates a block of memory of the given size:

malloc(size);

- Returns a pointer to the first byte of the allocated memory
 - Returns **NULL** if the memory could not be allocated
- You should assume the memory initially contains garbage
- **sizeof** is typically used to calculate the size you need

```
// allocate a 10-float array
float *arr = malloc(10 * sizeof(float));
if (arr == NULL)
    return 1;
arr[0] = 5.1;
```

Allocating memory: calloc

- Similar to malloc, but zeroes out allocated memory:
`calloc(nmemb, size);`
 - Returns an array of nmemb members, each of size bytes
 - Helpful for shaking out bugs
 - Slightly slower; preferred in non-performance-critical code

```
// allocate a 10-long array
long *arr = calloc(10, sizeof(long));
if (arr == NULL)
    return 1;
arr[0] = 5;
```

malloc vs. calloc

```
int malloc_calloc(void) {
    int i;
    unsigned char *mbuf, *cbuf;
    mbuf = malloc(5 * sizeof(char));
    cbuf = calloc(5, sizeof(char));
    for (i = 0; i < 5; i++) {
        printf("M[%d] = %02x, C[%d] = %02x\n",
               i, mbuf[i], i, cbuf[i]);
    }
    return 0;
}
```

M[0] = b8,	C[0] = 0
M[1] = 9d,	C[1] = 0
M[2] = 5e,	C[2] = 0
M[3] = 40,	C[3] = 0
M[4] = 4e,	C[4] = 0

Deallocating memory

- Releases the memory pointed to by ptr:

free(ptr);

- Must be the address of the first byte of allocated memory
 - I.e., the address that was returned from malloc or calloc
- This memory can then be allocated for something else later
- Good idea: set a pointer to NULL after you free it

```
// allocate a 10-long array
long *arr = calloc(10, sizeof(long));
// do something with arr...
free(arr);
arr = NULL;
```

The NULL pointer

- **NULL** refers to an address that is guaranteed to be invalid
 - On Linux/GCC, NULL is 0x0
 - Trying to dereference NULL causes a segmentation fault
- Good practice: set pointers to NULL whenever they *don't point to valid data*
 - Initialize to NULL
 - Set to NULL after free



Resizing allocated memory

- Resizes the allocated memory at `ptr` to `newsz`:

```
ptr = realloc(ptr, newsz);
```

- Note: this returns the new pointer!
 - If the memory is bigger than space allows, it has to be reallocated somewhere else, so the *address can change!*
 - **Always always always** get the return value from `realloc`!
- Behaves like `malloc` or `free` when convenient (see man page)

```
// allocate a 10-long array, then expand
long *arr = malloc(10 * sizeof(long));
arr[0] = 5;
arr = realloc(arr, 20 * sizeof(long));
arr[15] = -47;
```

Allocating structs

- You can malloc and free structs just like any other type
 - This is a very common scenario
 - sizeof is particularly helpful here

```
struct complex {
    double real;
    double imag;
};

// Allocates and initializes a new complex number
struct complex *cplx_alloc(double r, double i)
{
    struct complex *c = malloc(sizeof(struct complex));
    if (!c)
        return NULL;
    c->real = r;
    c->imag = i;
    return c;
}

// Provides a consistent way to free the struct
void cplx_free(struct complex *c)
{
    free(c);
}
```

Heap/stack in action

```
#include <stdlib.h>

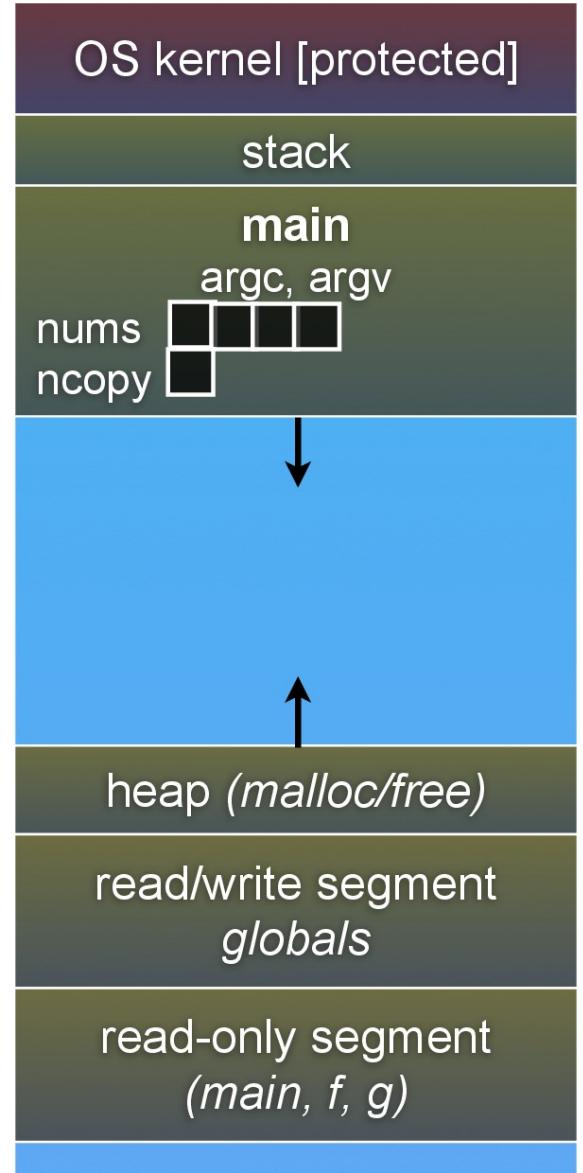
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    freencpy();
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

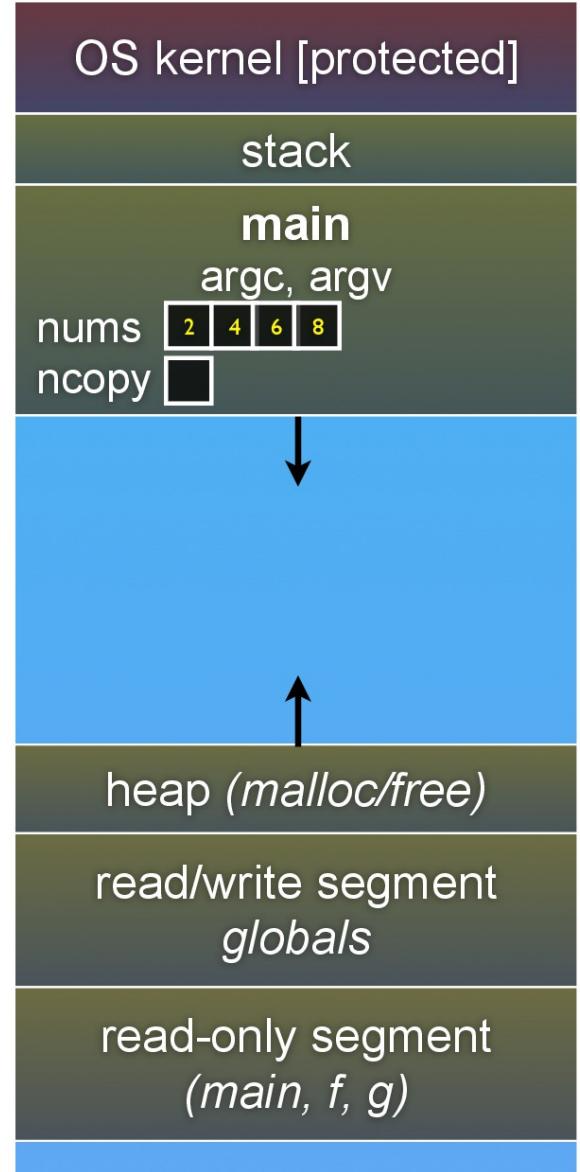
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

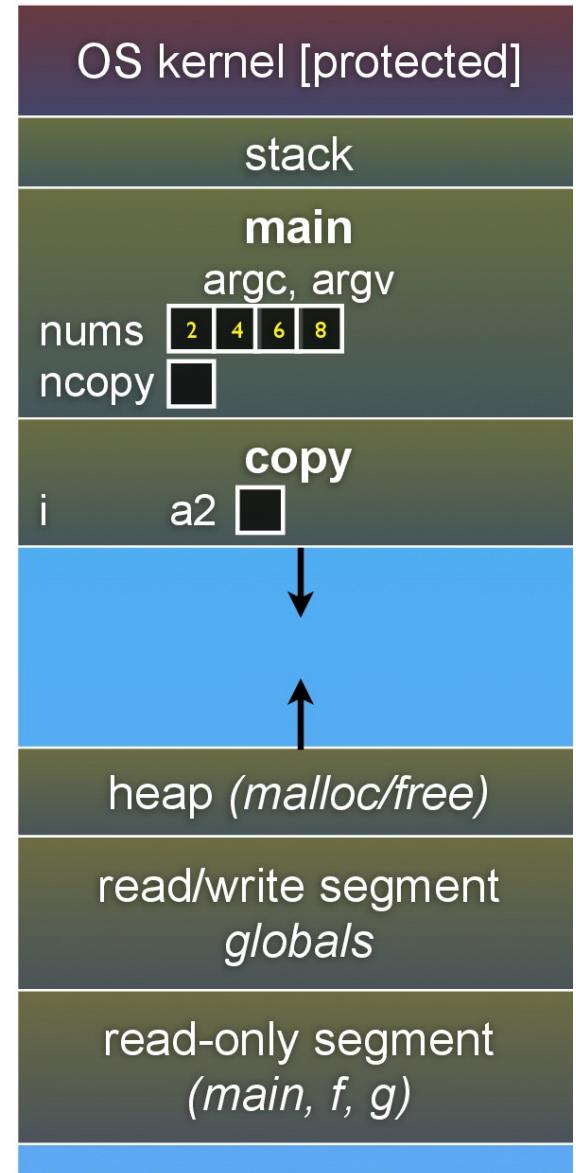
    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

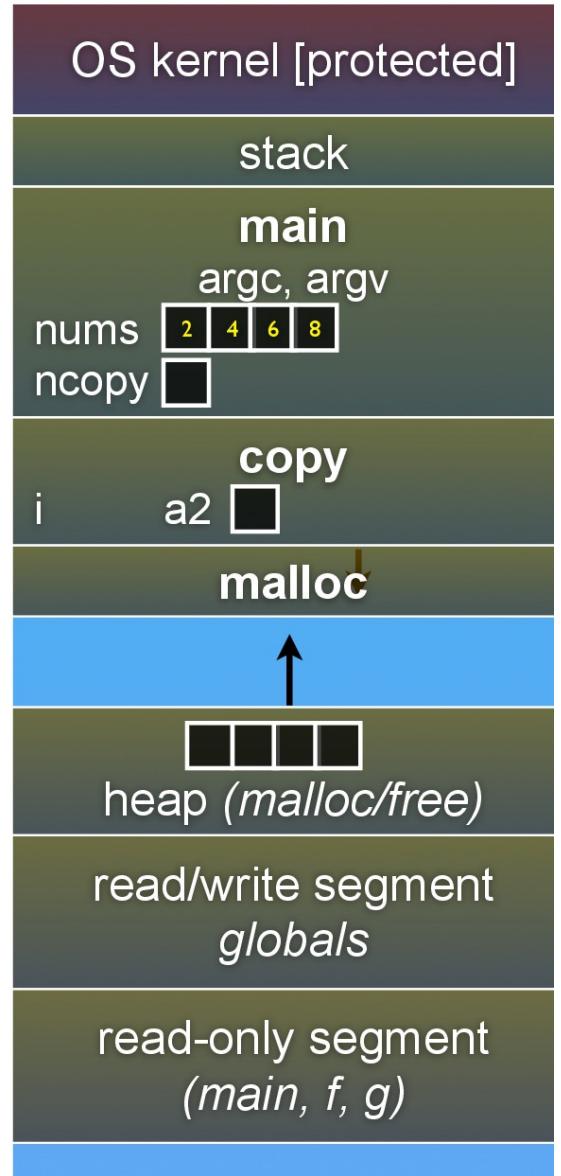
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

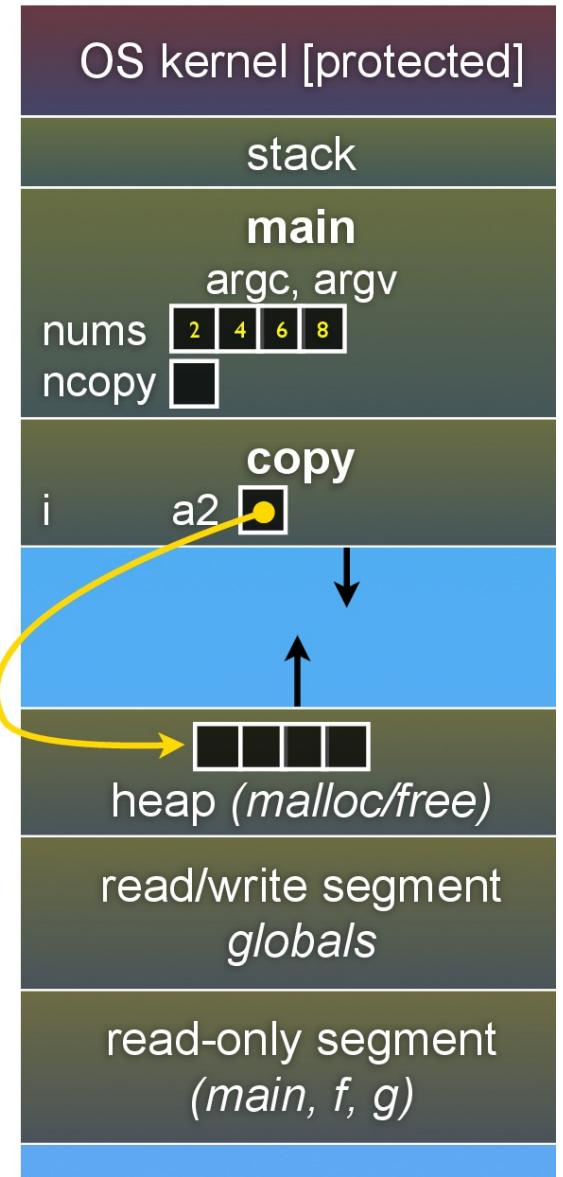
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

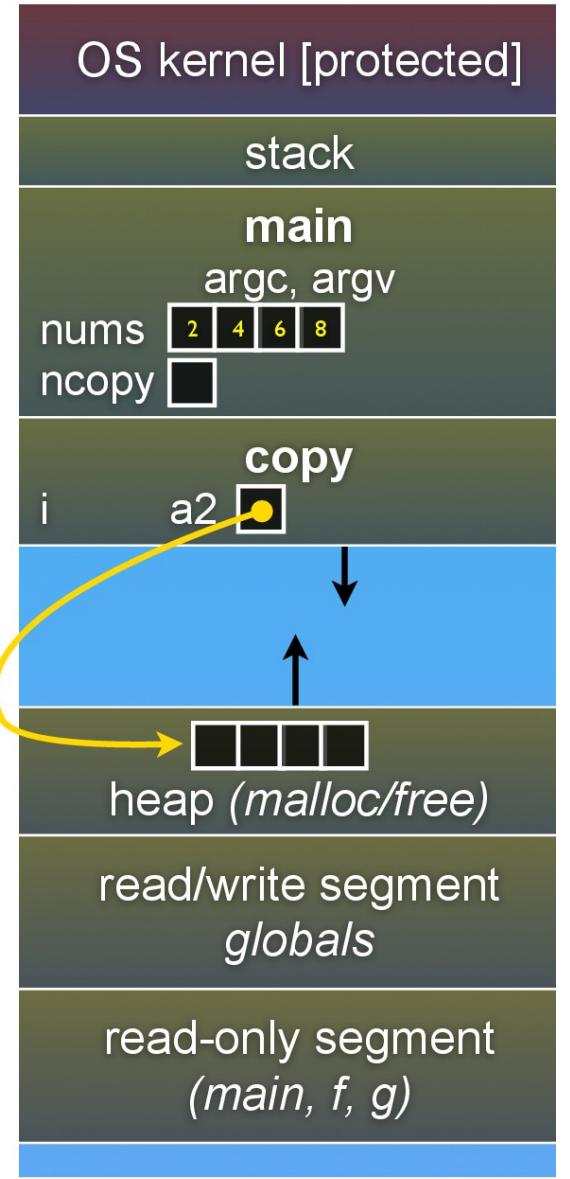
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

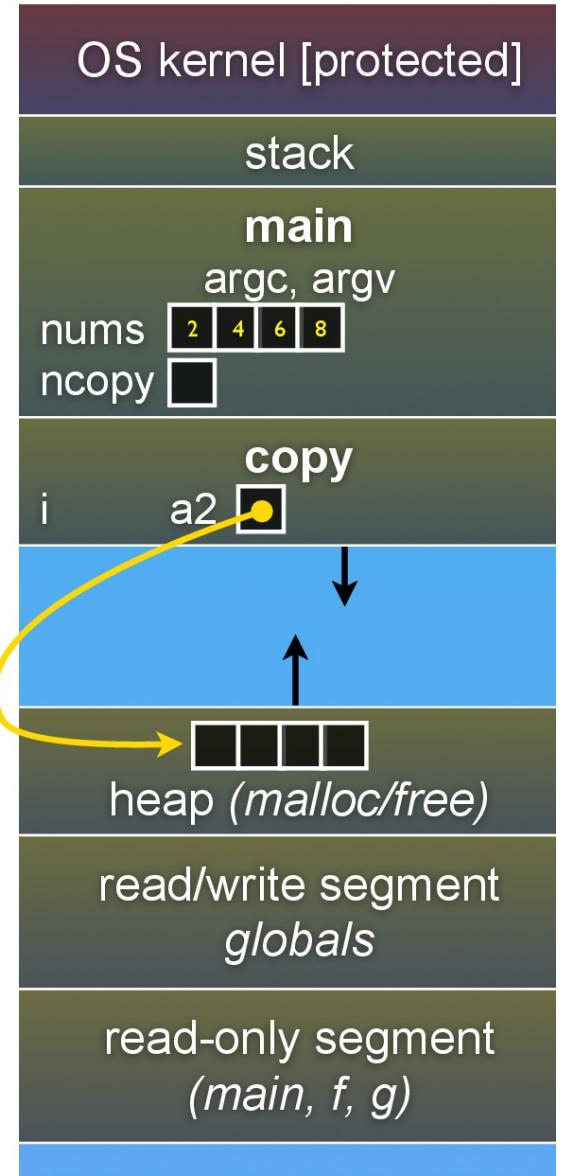
    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



arraycopy.c



Heap/stack in action



```
#include <stdlib.h>

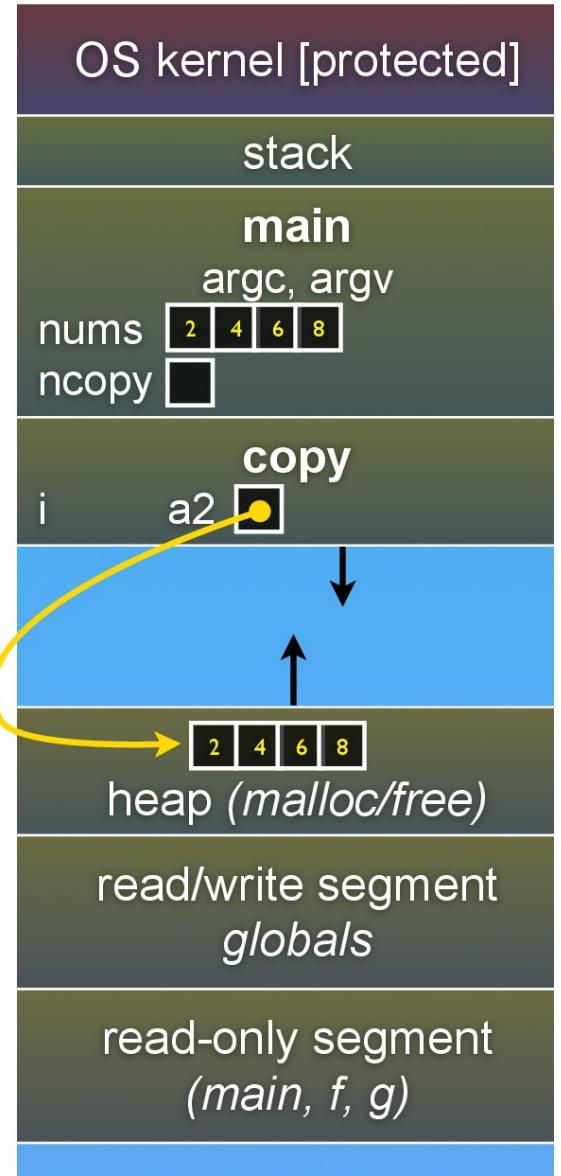
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

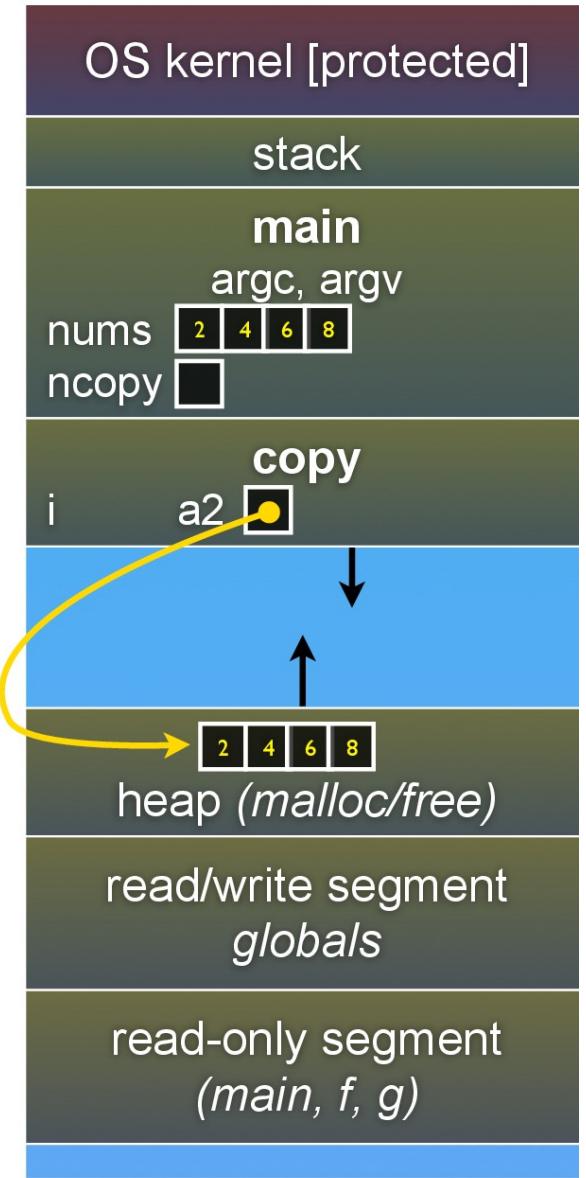
int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```

arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

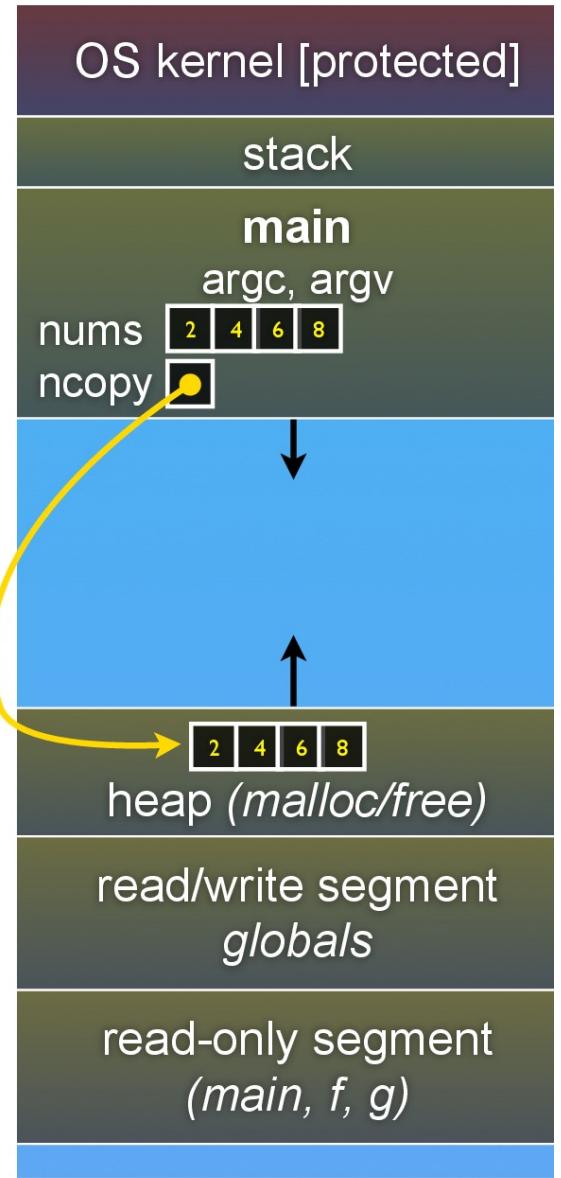
    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    freencpy;
    return 0;
}
```



arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

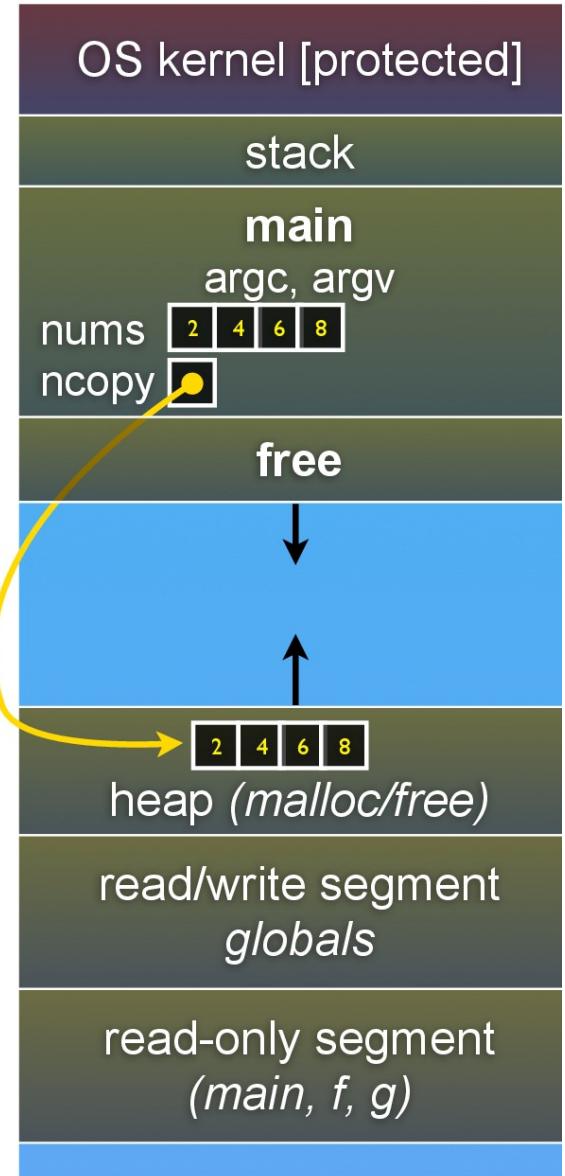
    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

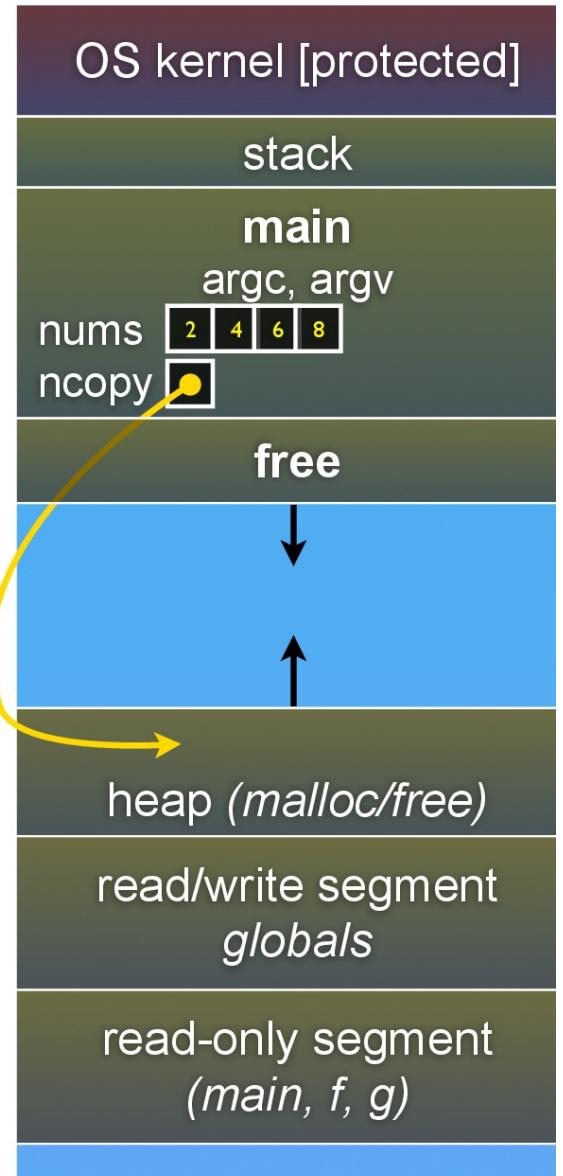
    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



arraycopy.c



Heap/stack in action

```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

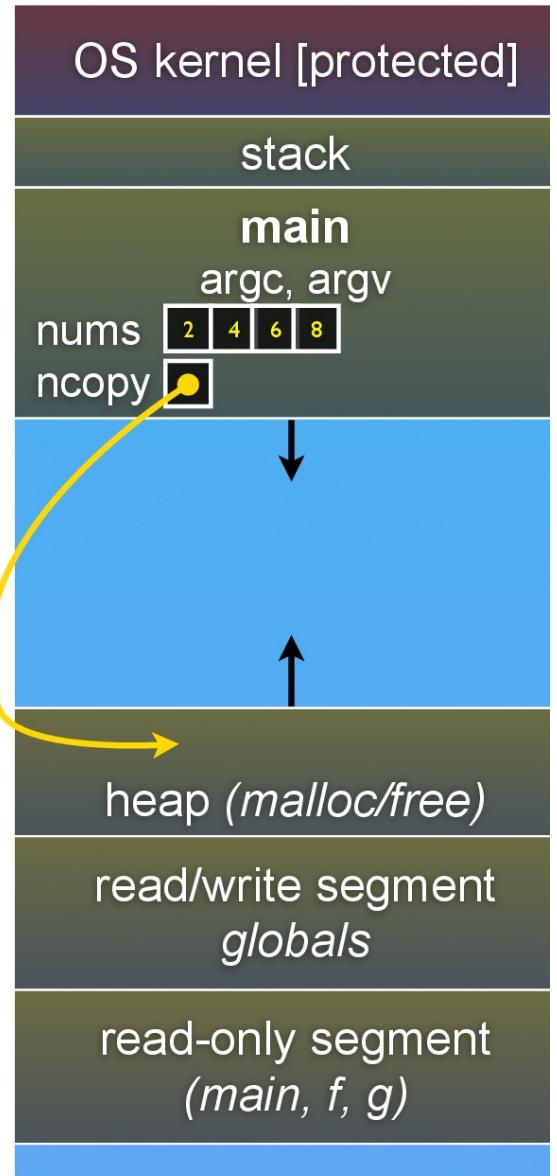
    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



arraycopy.c



Memory corruption

- There are all sorts of ways to use memory irresponsibly in C
 - Writing things where they don't belong or using allocation functions incorrectly can cause “memory corruption”

```
void use_memory_irresponsibly() {  
    int a[2];  
    int *b = malloc(2 * sizeof(int));  
    int *c;  
  
    // Many ways to corrupt memory...  
  
    a[2] = 5;          // assign past the end of an array  
    a[0] += 2;         // assume memory is zeroed by malloc  
    c = b + 3;         // mess up pointer arithmetic  
    free(&a[0]);       // free something that wasn't malloced  
    free(b);           // (this one is OK)  
    free(b);           // free something more than once (!)  
    b[0] = 5;          // use a pointer after it is freed  
    // ... and many more!  
}
```

Memory leak

- Happens when you allocate memory but forget to free it
 - Can be tricky to find
 - Sometimes hidden by a control statement like return, break, or continue
- Result: program's memory usage keeps growing
 - Slows down over time
 - Might be killed by the OS for using too much memory!

```
int string_length_leak() {  
    char *str = malloc(256);  
  
    scanf("%255s", str);  
    return strlen(str);  
}
```