



Systems and Internet Infrastructure Security

Institute for Networking and Security Research
Department of Computer Science and Engineering
Pennsylvania State University, University Park, PA



Memory Management (Part 3)

Devin J. Pohly <djpohly@cse.psu.edu>

Whence virtual memory?


- Every process begins with a certain amount of memory reserved for the heap
 - The top of the heap is known as the *program break*
 - Functions like `malloc` and `free` manage the heap by obtaining and releasing memory
 - You don't see it because it's handled for you



The program break

- The program break is moved up and down by `malloc` and `free`
 - Moved up to reserve more memory for the heap
 - Moved down to release some memory previously reserved for the heap

0xFFFFFFFF



0x00000000

brk and sbrk

- These functions are used to manage the program break
 - `int brk(void *addr)`: moves the program break to be at the address in `addr` (!)
 - This is an absolute address, so you can really mess things up – for example, by passing a pointer that points inside the stack.
 - `void *sbrk(long inc)`: moves the program break up by `inc` bytes
 - Positive `inc` increases the space reserved for the heap
 - Negative `inc` reduces the space reserved for the heap
 - Calling `sbrk(0)` will return the current program break without changing it

Let's try it

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void check_memory() {
    static void *last = NULL;

    void *ptr = sbrk(0);
    if (last)
        printf("The top of the heap is [%p] (%+d)\n", ptr, ptr - last);
    else
        printf("The top of the heap is [%p]\n", ptr);
    last = ptr;
}

int main()
{
    void *xptr[2048];
    int i;

    check_memory();
    xptr[0] = malloc(0x1000);
    check_memory();
    for (i = 1; i < 1024; i++) {
        xptr[i] = malloc(0x1000);
    }
    check_memory();
    for (i = 0; i < 1024; i++) {
        free(xptr[i]);
    }
    check_memory();
    return 0;
}
```

Let's try it

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void check_memory() {
    static void *last = NULL;

    void *ptr = sbrk(0);
    if (last)
        printf("The top of the heap is [%p] (%+d)\n", ptr, ptr - last);
    else
        printf("The top of the heap is [%p]\n", ptr);
    last = ptr;
}

int main()
{
    void *xptr[2048];
    int i;

    check_memory();
    xptr[0] = malloc(0x1000);
    check_memory();
    for (i = 1; i < 1024; i++) {
        xptr[i] = malloc(0x1000);
    }
    check_memory();
    for (i = 0; i < 1024; i++) {
        free(xptr[i]);
    }
    check_memory();
    return 0;
}
```

```
$ ./brktest
The top of the heap is [0x1b78000]
The top of the heap is [0x1b9a000] (+139264)
The top of the heap is [0x1f99000] (+4190208)
The top of the heap is [0x1b99000] (-4194304)
```


Some observations

- The program (via `malloc`) kept getting more and more memory over time.
- It released some, but not all of memory at the end.
 - Some “slack” is left over.

Top of heap	Change
0x1b78000	
0x1b9a000	+0x0022000
0x1f99000	+0x03ff000
0x1b99000	-0x0400000


An alternate method

- Some implementations of malloc don't use the program break; instead, they “**memory map**” a region to reserve more memory
 - In essence, this tells the OS that a certain range of memory should be available for use by the program.
 - This can be anywhere in the process address space (i.e., regardless of the program break).



Memory mapping

- In its most general use, the memory map system call “maps” (overlays) a file onto physical memory.
 - A program can then treat the file as random-access memory.
 - Very fast, because the OS “pages” the file content into memory in chunks (often 4kiB).
 - Note: memory mapping is often used to load program code when executing



Mapping memory

- To map memory, use the `mmap` function:

```
void *mmap(addr, len, prot, flags, fd, ofs);
```

- `void *addr`: address at which the file should be mapped
 - If this is `NULL`, the OS will choose where to map it.
- `size_t len`: how much memory to map
- `int prot`: protection bits (permissions)
- `int flags`: bits specifying the type of mapping
- `int fd`: file descriptor of the file to map (see I/O lectures)
- `off_t ofs`: starting offset within the file
- Returns a pointer to the region or “`(void *) -1`” on error

mmap protection bits

- The `prot` argument specifies the desired memory protection of the mapping.
 - It is either `PROT_NONE` or the bitwise OR of one or more of the following flags:
 - `PROT_READ`: memory may be read
 - `PROT_WRITE`: memory may be written
 - `PROT_EXEC`: memory may be executed
- These bits allow the program to police the use of mapped files/memory regions.

- Flags indicate what kind of mapping we are doing:
 - **MAP_SHARED**: ensures that updates to the mapping are written to the underlying file and visible to other processes that map it
 - **MAP_PRIVATE**: creates a private, copy-on-write mapping in which updates are not written to the underlying file or visible to other processes
- And one other option:
 - **MAP_ANONYMOUS**: the mapping is not backed by any file, and its contents are initialized to zero

Unmapping memory

- The complement to `mmap` is `munmap`:

```
void *munmap(addr, len);
```

 - `void *addr`: address of the memory to be released
 - Must be memory that was originally mapped with `mmap`!
 - `size_t len`: how much memory to release
- Note: any remaining changes to a `MAP_SHARED` mapping will be written to disk at this point

Mapping example

```
int mymap(char val) {
    // Local variables
    int fd, i;
    char *ptr;

    // Open the file and map 20 bytes to memory
    fd = open("mmap.dat", O_CREAT | O_RDWR);
    if (fd < 0) {
        perror("open");
        return 1;
    }

    ptr = mmap(NULL, 20, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    if (ptr == (void *) -1) {
        perror("mmap");
        return 1;
    }

    // Add the values
    for (i = 0; i < 20; i++) {
        ptr[i] = val;
    }

    // Release the mapped memory and close the file
    munmap(ptr, 20);
    close(fd);
    return 0;
}
```

Mapping example

```
int mymap(char val) {  
    // Local variables  
    int fd, i;  
    char *ptr;  
  
    // Open the file and map 20 bytes to memory  
    fd = open("mmap.dat", O_CREAT | O_RDWR);  
    if (fd < 0) {  
        perror("open");  
        return 1;  
    }  
  
    ptr = mmap(NULL, 20, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);  
    if (ptr == (void *) -1) {  
        perror("mmap");  
        return 1;  
    }  
}
```

```
(gdb) x/20xb ptr  
0x7ffff7ff6000: 0x7f    0x45    0x4c    0x46    0x02    0x01    0x01    0x00  
0x7ffff7ff6008: 0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00  
0x7ffff7ff6010: 0x02    0x00    0x3e    0x00
```

```
(gdb) next
```

```
...
```

```
(gdb) x/20xb ptr  
0x7ffff7ff6000: 0x58    0x58    0x58    0x58    0x58    0x58    0x58    0x58  
0x7ffff7ff6008: 0x58    0x58    0x58    0x58    0x58    0x58    0x58    0x58  
0x7ffff7ff6010: 0x58    0x58    0x58    0x58
```

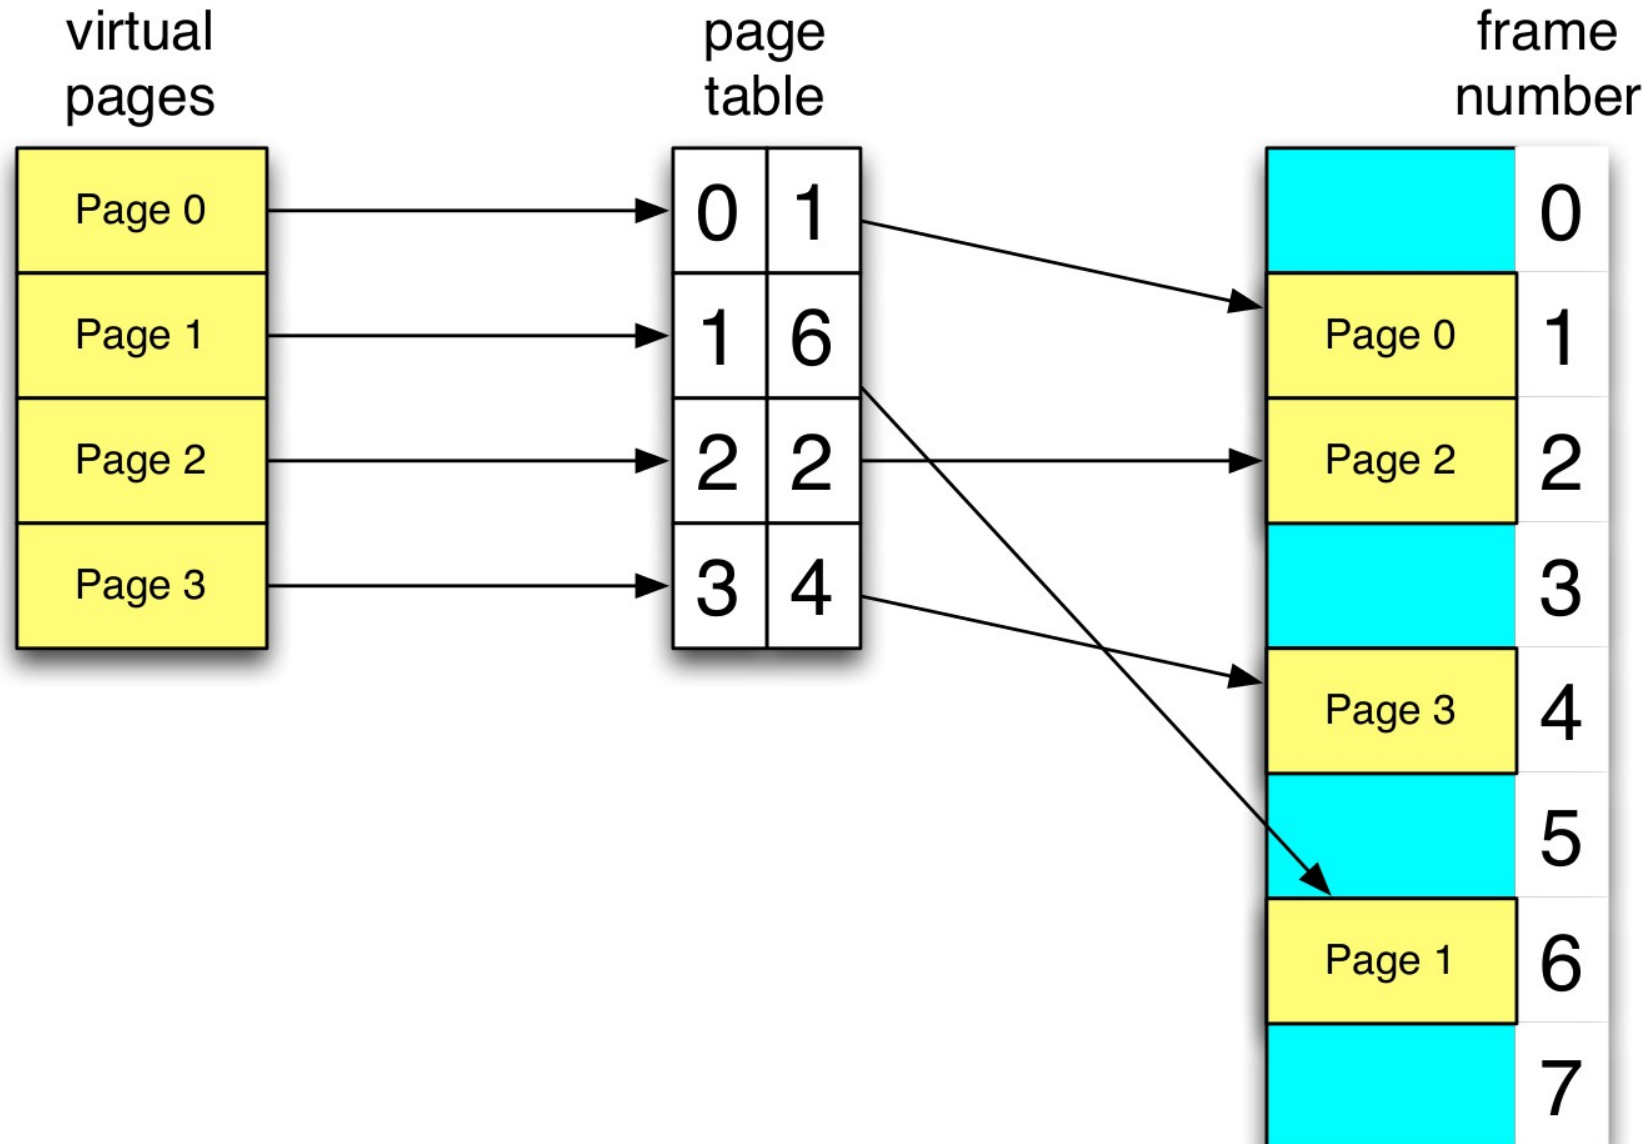

An alternate method

- So an implementation of malloc and free can use:
 - `mmap` instead of `brk/sbrk`
 - `munmap` instead of `brk/sbrk`
- This is a design choice made when implementing the standard library
 - e.g., the GNU C library (glibc)



- Programs have a virtual address space (say, 1 MiB)
- OS must fetch data from either physical memory or disk
 - Done by a mechanism called *paging* (or *demand paging*)
 - Divide the virtual address space into units called *pages*, each of which is a fixed size (often 4kiB)
 - For example 1 MiB virtual address space has 256 4kiB pages.
 - Physical memory is divided into *frames*, usually the same size as pages
 - For example, we may have only 32 4kiB frames.
- OS tracks where pages are stored in physical memory
 - Maintained in a data structure called the *page table*
 - Maps virtual addresses to physical addresses

Example page table



Multiple processes

