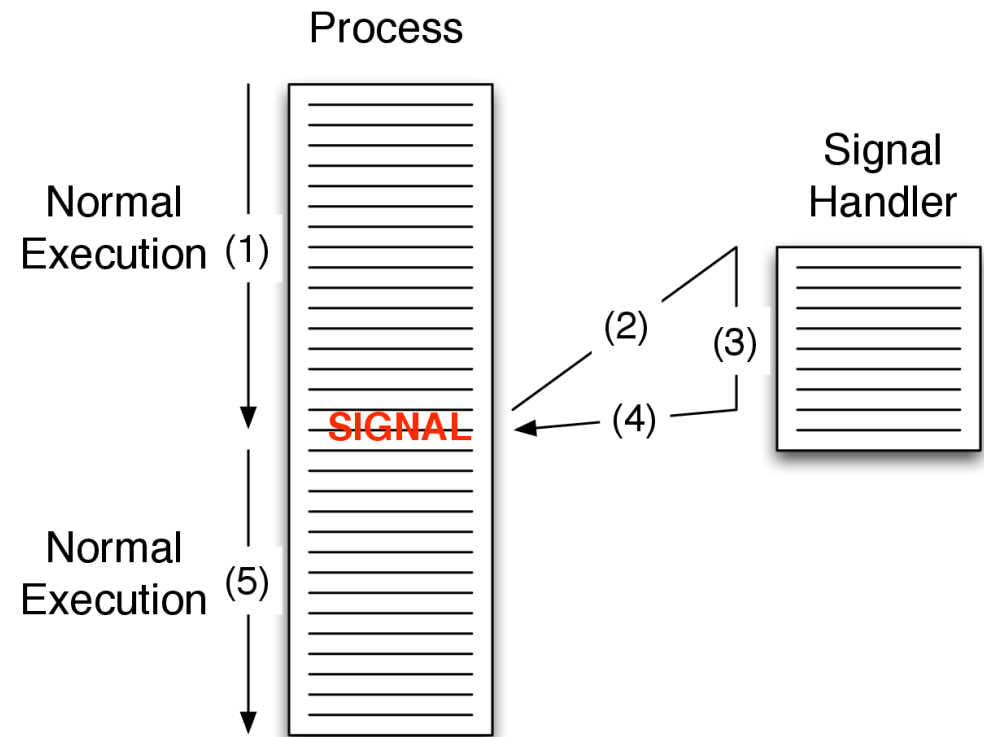


Signals

Devin J. Pohly <djpohly@cse.psu.edu>

Unix signals

- *Signals* are special messages sent through the OS to tell a process (or thread) of some request or event
- Process execution stops and special *signal handler* code runs
 - The process may resume operation after the signal handler finishes
- You can only send signals to your own processes
 - The OS and root can send signals to any process



Signal names

- Each signal is identified by a number:

```
/* Signals. */
#define SIGHUP      1    /* Hangup (POSIX).  */
#define SIGINT      2    /* Interrupt (ANSI). */
#define SIGQUIT     3    /* Quit (POSIX).   */
#define SIGABRT     6    /* Abort (ANSI).   */
#define SIGFPE      8    /* Floating-point exception (ANSI). */
#define SIGKILL     9    /* Kill, unblockable (POSIX). */
#define SIGSEGV    11    /* Segmentation violation (ANSI). */
#define SIGPIPE    13    /* Broken pipe (POSIX). */
#define SIGALRM    14    /* Alarm clock (POSIX). */
#define SIGTERM    15    /* Termination (ANSI). */
#define SIGCHLD    17    /* Child status has changed (POSIX). */
#define SIGCONT    18    /* Continue (POSIX). */
#define SIGSTOP    19    /* Stop, unblockable (POSIX). */
#define SIGSYS     31    /* Bad system call. */
```

- All the signals are #defined in header files:
 - To use them, `#include <signal.h>`
 - Actual definitions are in `/usr/include/bits/signum.h`
 - Note: use names when possible; the numbers are not portable!

Signals as process control

- The operating system uses signals to control how the process runs, or stops running.

- Signals are sent on errors:

```
#define SIGILL      4    /* Illegal instruction (ANSI).  */
#define SIGTRAP     5    /* Trace trap (POSIX).  */
#define SIGABRT     6    /* Abort (ANSI).  */
#define SIGBUS      7    /* BUS error (4.2 BSD).  */
#define SIGFPE      8    /* Floating-point exception (ANSI).  */
#define SIGSEGV     11   /* Segmentation violation (ANSI).  */
```

- Signals can be used by applications as well:

```
#define SIGUSR1     10   /* User-defined signal 1 (POSIX).  */
#define SIGUSR2     12   /* User-defined signal 2 (POSIX).  */
```

- Signals can control process execution:

```
#define SIGINT      2    /* Interrupt (ANSI).  */
#define SIGKILL     9    /* Kill, unblockable (POSIX).  */
#define SIGTERM     15   /* Termination (ANSI).  */
#define SIGCONT     18   /* Continue (POSIX).  */
#define SIGSTOP     19   /* Stop, unblockable (POSIX).  */
```

Process IDs

- Every process running on the OS is given a unique number called its *process ID* (or *pid*)
 - This is what is used in the OS and for process control to reference one specific running program instance
- To see process IDs for current processes, use the *ps* utility

```
$ ps -U djpohly
PID TTY          TIME CMD
 699 ?            00:00:00 systemd
 714 ?            00:00:00 dbus-daemon
 729 ?            01:36:18 mpd
1021 ?            00:00:00 dbus-daemon
1022 ?            16:07:09 jackd
2335 pts/7        00:00:00 bash
2684 pts/7        00:00:00 ps
2685 pts/7        00:00:00 xsel
3710 ?            00:00:00 bash
3711 ?            00:00:00 startx
3730 ?            00:00:00 xinit
3731 ?            01:44:09 X
3734 ?            00:00:00 .xinitrc
3739 ?            00:06:27 xcompmgr
3741 ?            00:03:32 urxvtd
3745 ?            00:00:12 dwm
8254 ?            01:09:07 firefox
8660 pts/3        00:00:00 bash
8718 pts/4        00:00:00 bash
9378 ?            00:00:00 dbus-launch
9379 ?            00:00:00 dbus-daemon
9381 ?            00:00:00 at-spi-bus-laun
9387 ?            00:00:09 at-spi2-registr
10470 pts/0        00:00:05 mutt
12920 pts/1        00:00:00 dtach
12922 ?            00:00:00 dtach
12923 pts/2        00:00:20 mutt
16119 ?            00:00:00 mpc
23895 ?            00:00:00 mpdfilter
26870 ?            00:00:03 zathura
27372 ?            00:00:00 loimpress
27388 ?            00:00:00 oosplash
27400 ?            00:00:47 soffice.bin
...
```

The `kill` command

- To send a signal to a process, use the `kill` command (named after its most common use...):

`kill -SIG pid(s)`

- **SIG**: the signal name (e.g., `ABRT`) or number (e.g., `6`)
 - `SIGTERM` is the default if no `-SIG` argument is given
- **pid(s)**: process IDs of one or more processes to which to send the given signal

```
$ ./signals
Main screen turn on...
We get SIGHUP!
Main screen turn on...
^CWe get SIGINT!
Main screen turn on...
We get SIGTERM!
Exiting cleanly
```

```
$ ps -U djpohly | grep signals
11921 pts/7      00:00:00 signals
$ kill -1 11921
$ # typed Ctrl-C on program
$ kill -TERM 11921
```

SIGTERM vs. SIGKILL

- **SIGTERM** interrupts the program and asks it nicely to shut down
 - Sometimes this will not work (e.g., if the process is in a locked state)
 - It is often desirable to add a signal handler for SIGTERM so that your program can clean up memory, close files, and terminate of its own accord
 - This is referred to as a *graceful shutdown*
- **SIGKILL** kills the process unceremoniously
 - Can lead to inconsistent state, buffers not written to file, certain system resources not being released, etc.
 - Effective, but should be a last resort
 - If someone tells you to “**kill -9**” a process, that’s SIGKILL

The `killall` command

- The `killall` program sends signals to *all* instances of a particular program:

`killall -SIG name(s)`

- `SIG`: same as for `kill`
- `name(s)`: one or more program names to which to send the signals

```
$ ./signals
Main screen turn on...
We get SIGHUP!
Main screen turn on...
We get SIGINT!
Main screen turn on...
Killed
```

```
$ killall -HUP signals
$ killall -2 signals
$ killall -KILL signals
$ # ^ that was overkill :(
```


Signal dispositions

- If a process receives a signal for which it does not provide a handler, the signal's default action or *disposition* is invoked:
 - Most signals: terminate the process
 - Some signals: terminate the process and generate a core dump for debugging if enabled (e.g., SIGSEGV, SIGABRT)
 - A couple of signals: ignored
 - Certain signals: pause the process (SIGSTOP) or resume execution (SIGCONT)

Meanwhile, back in C...

- The `kill` function sends a signal to a process:

```
int kill(pid_t pid, int sig)
```

- `pid`: process ID
- `sig`: signal number (use the named constants, e.g., `SIGTERM`)
- Return value: 0 for success



Raising a signal

- *Raising* a signal simply means a process sending the signal to itself:

```
int raise(int sig);
```

- There are a range of reasons why a process might want to do this:
 - Suspend itself (SIGSTOP)
 - Terminate itself or its threads (SIGTERM, SIGKILL)
 - User-defined signals (SIGUSR1, ...)
- Return value: 0 for success

Process-defined handlers

- You can create your own signal handler simply by creating a function with the following signature:

`void funcname(int argname)`

- To set this function as the handler for a particular signal, you get a *function pointer* to the handler and pass it to the `signal` function:

`sighandler_t signal(int signum, sighandler_t handler)`

- “`sighandler_t`” is a function pointer typedef
- Return value: previous signal handler, or `SIG_ERR`

```
void signal_handler(int num) {  
    printf("Got signal %d\n", num);  
}
```

```
signal(SIGHUP, signal_handler);  
signal(SIGINT, signal_handler);
```

Function pointers

- A *function pointer* is a variable that points to a function; it can be assigned, passed as a parameter, and called like a regular function.
- To declare a function pointer:
*rettype (*var)(argtypes);*
 - *rettype*: return type of the function
 - *var*: name of the pointer variable being declared
 - *argtypes*: types of the arguments, separated by commas
- Probably the most confusing syntax in all of C/C++!
 - Frequently a typedef (like `sighandler_t`) is used to help

Using function pointers

- Just like the name of an array is a pointer to the array, the name of a function is a pointer to the function
 - C will let you use the & operator with the function name for clarity, but it's not required

```
void foo(int num) {  
    printf("foo: num=%d\n", num);  
}  
  
void bar(int num) {  
    printf("bar: num=%d\n", num);  
}  
  
int main(int argc, char **argv) {  
    // Declare, assign, and call fp  
    void (*fp)(int);  
    fp = foo;  
    fp(42);  
    fp = &bar;  
    fp(74);  
    return 0;  
}
```

```
$ ./fptest  
foo: num=42  
bar: num=74
```

A new approach

- The `sigaction` function changes the action taken by a process on receipt of a specific signal:

```
int sigaction(int signum, const struct sigaction *act,  
              struct sigaction *oldact)
```

- `signum`: what signal to handle
- `act`: pointer to a structure containing information about the new handler, `NULL` means don't change the existing one
- `oldact`: pointer to a structure where the old handler info will be stored, `NULL` means we don't want it
- `man sigaction` to see what the `sigaction` struct contains

Why another API?

- `sigaction` is newer, cleaner, and generally preferred:
 - `signal` does not block other signals from arriving during the current handler (which can cause a race condition); `sigaction` can block other signals until the current handler returns.
 - On some systems, `signal` resets the signal to its default disposition after the handler executes. On others, it doesn't.
 - `sigaction` provides flags for more control over signal behavior:
 - `SA_NODEFER`: don't suspend signals while handler is executing
 - `SA_ONSTACK`: use an alternate stack for the signal handler
 - `SA_RESETHAND`: restore the signal to default disposition when the handler is executed
 - And others, see `man sigaction`

All at once now

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

int running = 1;

void signal_handler(int num)
{
    switch (num) {
        case SIGHUP:
            printf("We get SIGHUP!\n");
            break;
        case SIGINT:
            printf("We get SIGINT!\n");
            break;
        case SIGTERM:
            // End the program cleanly
            printf("We get SIGTERM!\n");
            running = 0;
            break;
    }
}
```

```
int main(int argc, char **argv)
{
    struct sigaction new_act, old_act;
    new_act.sa_handler = signal_handler;
    new_act.sa_flags = 0;
    sigaction(SIGINT, &new_act, &old_act);

    signal(SIGHUP, signal_handler);
    signal(SIGTERM, signal_handler);

    while (running) {
        printf("Main screen turn on...\n");
        // Wait for a signal
        pause();
    }

    // Do code cleanup here
    printf("Exiting cleanly\n");

    return 0;
}
```