



Systems and Internet Infrastructure Security

Network and Security Research Center
Department of Computer Science and Engineering
Pennsylvania State University, University Park PA

CMPSC 311 - Introduction to Systems Programming Module: Concurrency

Professor Patrick McDaniel
Fall 2013


Sequential Programming

- Processing a network connection as it arrives and fulfilling the exchange completely is sequential processing
 - ▶ i.e., connections are processed in *sequence* of arrival

```
listen_fd = Listen(port);  
while(1) {  
    client_fd = accept(listen_fd);  
    buf = read(client_fd);  
    write(client_fd, buf);  
    close(client_fd);  
}
```


Whither sequential?

- Benefits
 - super simple to build
 - very little persistent state to maintain
- Disadvantages
 - incredibly poorly performing
 - one slow client causes *all* others to block
 - poor utilization of network, CPU
 - cannot interleave processing



Whither sequential?

- Benefits
 - super simple to build
 - very little persistent state to maintain
- Disadvantages
 - incredibly poorly performing
 - one slow client causes *all* others to block
 - poor utilization of network, CPU
 - cannot interleave processing



Think about this way: if the class took the final exam sequentially, it would take 6 days and 16 hours to complete the exam!

An alternate design ...

- Why not process multiple requests at the same time, interleaving processing while waiting for other dependent actions (DB dips) to complete?
- This is known as *concurrent processing* ...
 - Process multiple requests concurrently
- We have/will explore three concurrency approaches
 1. Event programming using `select()` [see last lecture]
 2. Creating “child” processes using `fork()`
 3. Creating multithreaded programs using `pthread`s

Concurrency with processes


- The server process blocks on `accept()`, waiting for a new client to connect
 - ▶ when a new connection arrives, the parent calls `fork()` to create a another process
 - ▶ the child process handles that new connection, and `exit()`'s when the connection terminates
- Children become “zombies” after death
 - ▶ `Wait()` to “reap” children




Graphically

● server


Graphically




Graphically




Graphically



Graphically



Graphically




child `exit()`'s / parent `wait()`'s


Graphically





Graphically



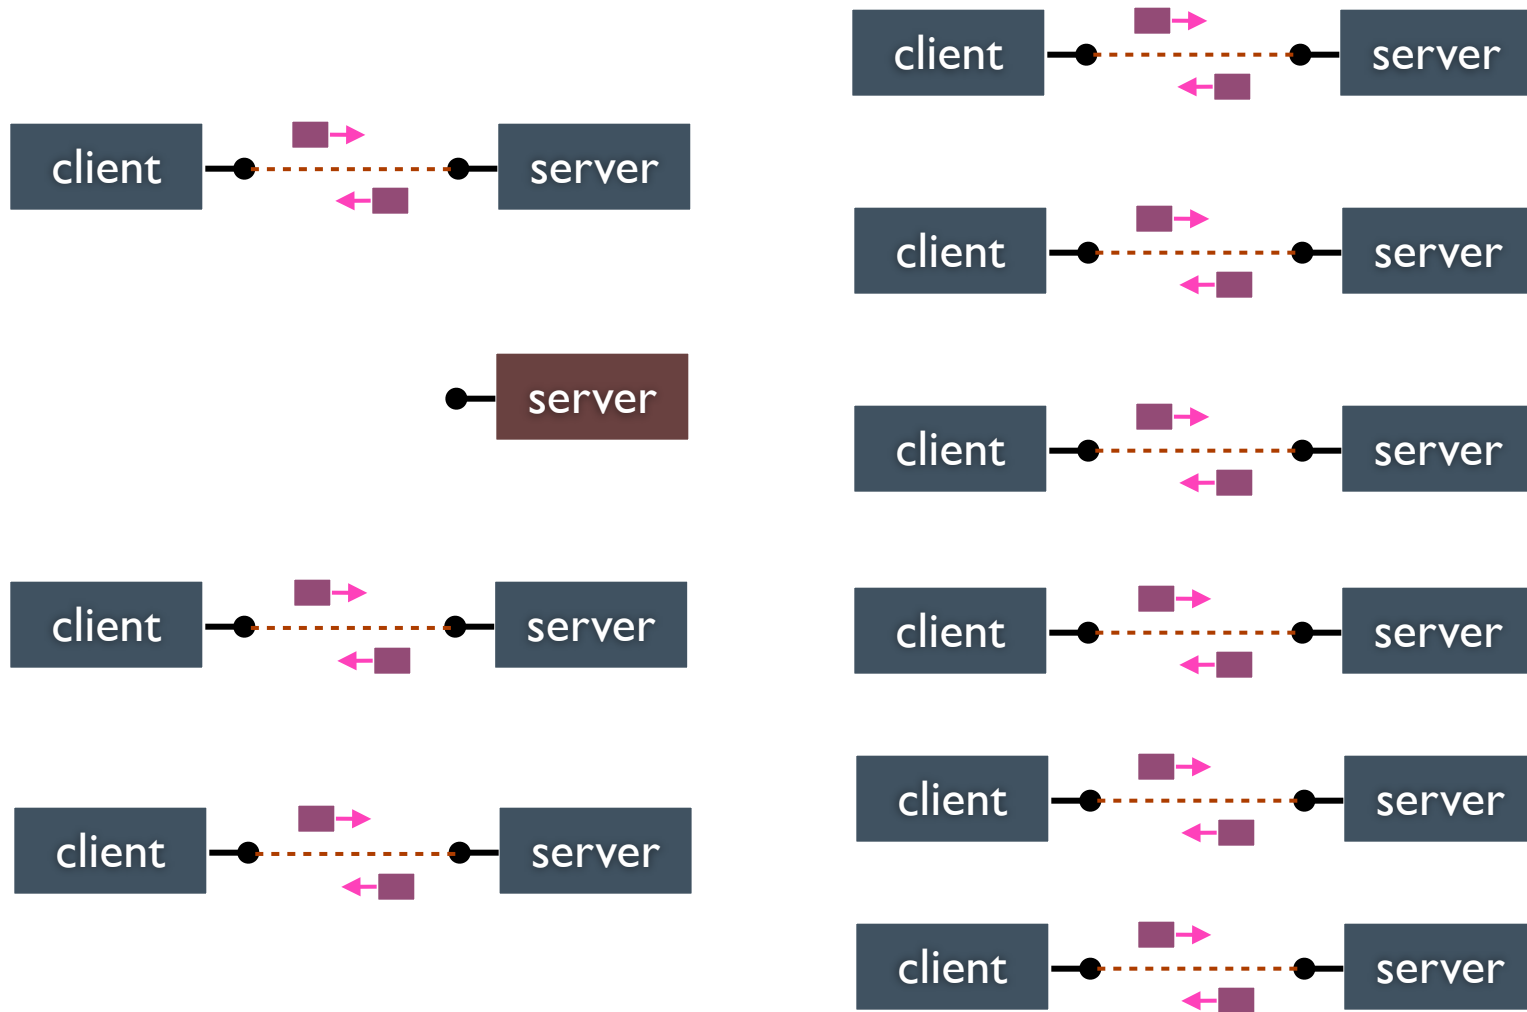
Graphically



Graphically



Graphically



fork()

- The `fork` function creates a new process by duplicating the calling process.

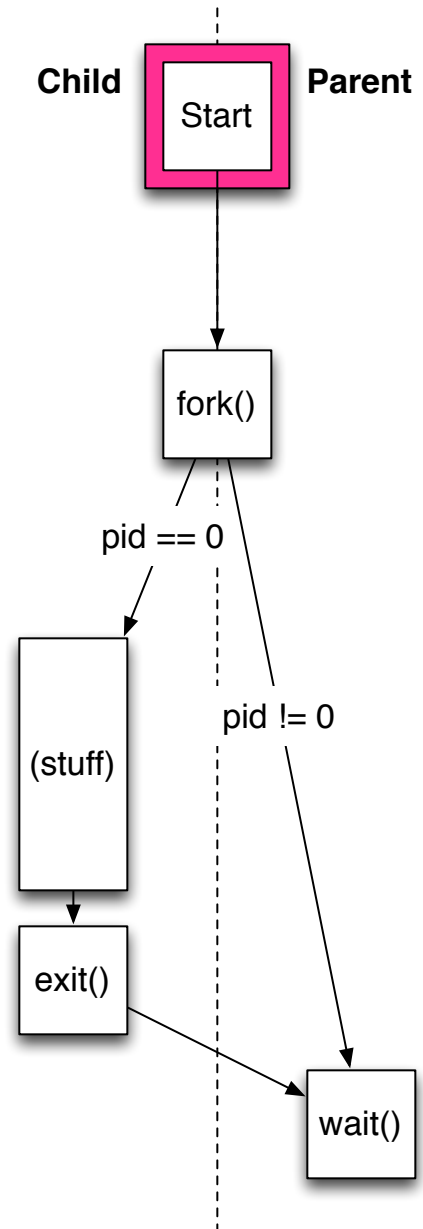
```
pid_t fork(void);
```

- The new `child` process is an exact duplicate of the calling `parent` process, except that it has its own process ID and pending signal queue
- The `fork()` function returns
 - ▶ 0 (zero) for the child process OR ...
 - ▶ The child's `process ID` in the parent code

Idea: think about duplicating the process state and running

Process control

- Parent
 - fork (pid == child PID)
 - wait for child to complete (maybe)
- Child
 - begins at fork (pid == 0)
 - runs until done
 - calls exit



exit()

- The `exit` causes normal process termination with a return value

```
void exit(int status);
```






- Where
 - `status` is sent to the to the parent (&0377)
- Note: `exit` vs. `return` from a C program
 - `return` simply returns a normal execution
 - `exit` kills the process via system call
 - Consequence: `return` is often cleaner for certain things that run `atexit` (more of a C++ concern)

wait()

- The `wait` function is used to wait for state changes in a child of the calling process (e.g., to terminate)

```
pid_t wait(int *status);
```

- Where
 - ▶ returns the process ID of the child process
 - ▶ `status` is return value set by the child process



wait () and the watchdog

- Often you will want to safeguard a process by “watching” it execute. If the process terminates then you take some action (when `wait ()` returns).
 - ▶ This is called a *watchdog process* (or watchdog fork)
 - ▶ Idea: you fork the process and simply wait for it to terminate
 - Take some action based on the idea that the termination should not have happened or that some recovery is needed.
 - ▶ e.g., restart a web server when it crashes
- Watchdogs are very often used for critical services.

Putting it all together ...

```
int main( void ) {

    pid_t childpid; // Child process ID
    int status;      // Childs exist status

    printf( "Stating our process control example ...\n" );
    childpid = fork(); // Fork the process

    if ( childpid >= 0 ) {
        if ( childpid == 0 ) { // In child process
            printf( "Child processing,\n" );
            sleep( 1 );
            exit( 19 );
        } else { // In parent process
            printf( "Parent processing (child=%d),\n", childpid );
            wait( &status );
            printf( "Child exited with status=%d.\n", WEXITSTATUS(status) );
        }
    }
    else {
        perror( "What the fork happened" );
        exit( -1 );
    }
    return( 0 );
}
```

Putting it all together ...

```
int main( void ) {

    pid_t childpid; // Child process ID
    int status;      // Childs exist status

    printf( "Stating our process control example ...\n" );
    childpid = fork(); // Fork the process

    if ( childpid >= 0 ) {
        if ( childpid == 0 ) { // In child process
            printf( "Child processing,\n" );
            sleep( 1 );
            exit( 19 );
        } else { // In parent process
            printf( "Parent processing (child=%d),\n", childpid );
            wait( &status );
            printf( "Child exited with status=%d.\n", WEXITSTATUS(status) );
        }
    }
    else {
        perror( "What the fork happened" );
        exit( -1 );
    }
    return( 0 );
}
```

```
$ ./forker
Stating our process control example ...
Parent processing (child=3530),
Child processing.
Child exited with status=19.
$
```


Process Control Tradeoffs

- Benefits
 - ▶ almost as simple as sequential
 - ▶ in fact, most of the code is identical!
 - ▶ parallel execution; good CPU, network utilization
 - ▶ often better security (isolation)
- Disadvantages
 - ▶ processes are heavyweight
 - ▶ relatively slow to fork
 - ▶ context switching latency is high
 - ▶ communication between processes is complicated

Concurrency with threads

- A single process handles all of the connections
 - ▶ but, a *parent thread* forks (or dispatches) a new thread to handle each connection
 - ▶ the *child thread*:
 - handles the new connection
 - exits when the connection terminates

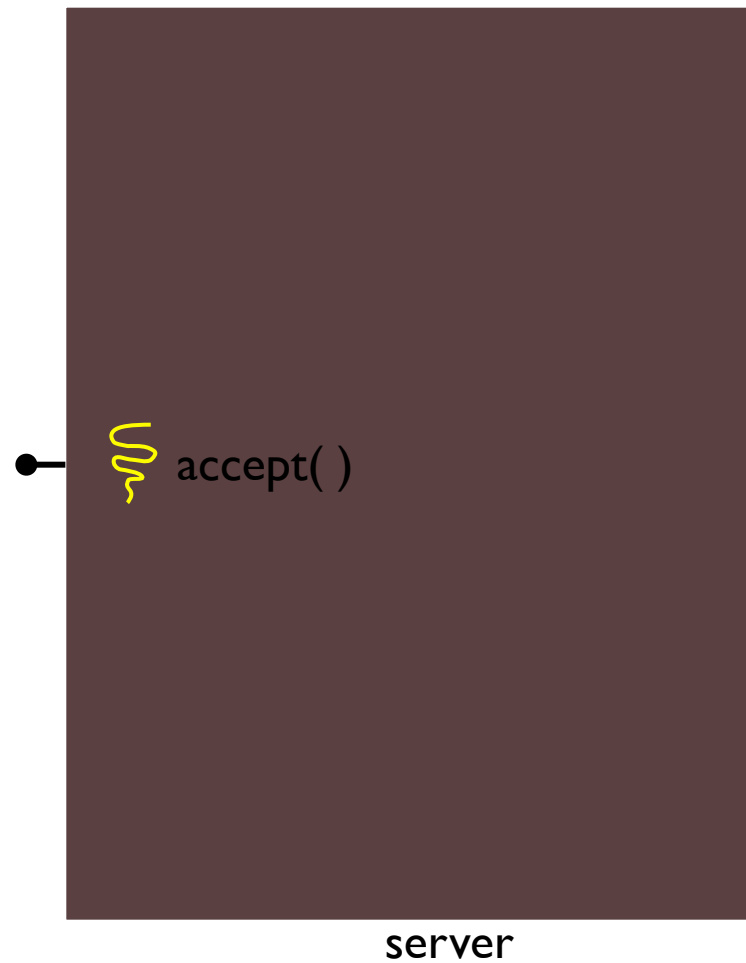


Note: you can create as many threads as you want (up to a system limit)


- A **thread** is defined as an independent stream of instructions that can be scheduled to run as such by the operating system.
 - ▶ To the software developer, the concept of a "procedure" that runs independently from its main program.
 - ▶ To go one step further, imagine a main program that contains a number of procedures. Now imagine all of these procedures being able to be scheduled to run simultaneously and/or independently by the operating system. That would describe a "multi-threaded" program.

Idea: “forking” multiple threads of execution in one process!


Threads (cont.)




Threads (cont.)




Threads (cont.)




Threads (cont.)



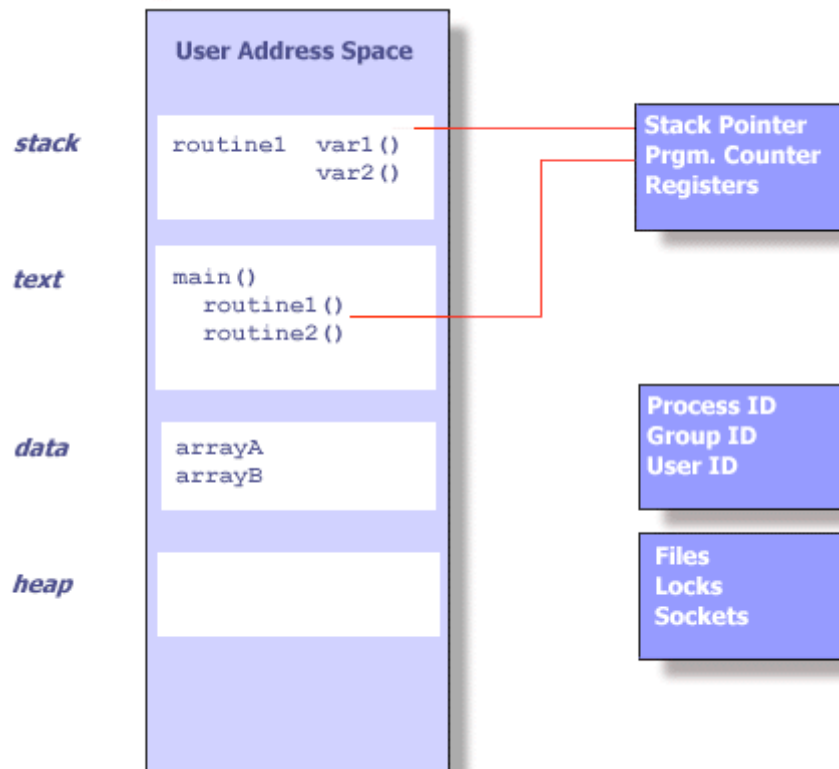
Threads (cont.)



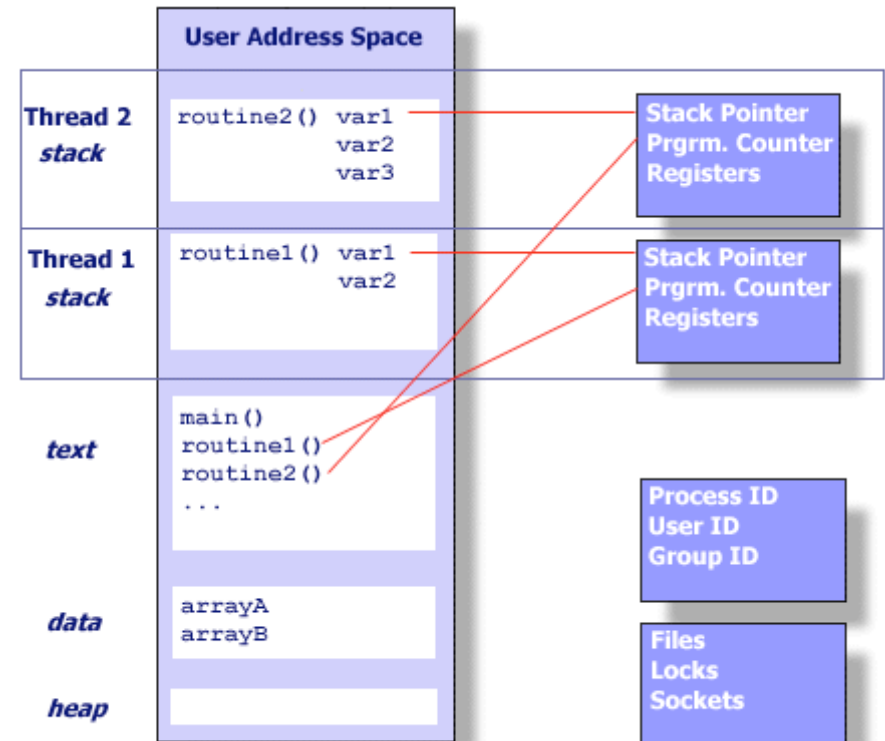
Threads (cont.)



Threads (cont.)



UNIX Process



... and with threads

Threads (cont.)

- This independent flow of control is accomplished because a thread maintains its own:
 - ▶ Stack pointer
 - ▶ Registers
 - ▶ Scheduling properties (such as policy or priority)
 - ▶ Set of pending and blocked signals
 - ▶ Thread specific data.


Thread Summary

- Exists within a process and uses the process resources
- Has its own independent flow of control as long as its parent process exists and the OS supports it
- Duplicates only the essential resources it needs to be independently “schedulable”
- May share the process resources with other threads that act equally independently (and dependently)
- Dies if the parent process dies - or something similar
- Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.

- Because threads within the same process share resources:
 - ▶ Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads.
 - ▶ Two pointers having the same value point to the same data.
 - ▶ Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer.

Thread control

- main
 - ▶ `pthread_create()` (create thread)
 - ▶ wait for thread to finish via `pthread_join()` (maybe)
- thread
 - ▶ begins at function pointer
 - ▶ runs until `pthread_exit()`



pthread_create()

- The pthread_create function starts a new thread in the calling process.

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine) (void *),  
                  void *arg);
```

- Where,
 - ▶ `thread` is a pthread library structure holding thread info
 - ▶ `attr` is a set of attributes to apply to the thread
 - ▶ `start_routine` is the thread function pointer
 - ▶ `arg` is an opaque data pointer to pass to thread

pthread_join()

- The pthread_join function waits for the thread specified by thread to terminate.

```
int pthread_join(pthread_t thread, void **retval);
```


- Where,
 - ▶ `thread` is a pthread library structure holding thread info
 - ▶ `retval` is a double pointer return value

pthread_exit()

- The pthread_exit function terminates the calling thread and returns a value

```
void pthread_exit(void *retval);
```

- Where,
 - `retval` is a pointer to a return value
 - **Note:** better be dynamically allocated because the thread stack will go away when the thread exits



Putting it all together ...

```
typedef struct {
    int num;
    const char *str;
} MY_STRUCT;

void * thread_function( void * arg ) {
    MY_STRUCT *val = (MY_STRUCT *)arg; // Cast to expected type
    printf( "Thread %lx has vaules %x,%s]\n", pthread_self(), val->num, val->str );
    pthread_exit( &val->num );
}

int main( void ) {
    MY_STRUCT v1 = { 0x12345, "Val 1" };
    MY_STRUCT v2 = { 0x54312, "Val 2" };
    pthread_t t1, t2;
    printf( "Starting threads\n" );
    pthread_create( &t1, NULL, thread_function, (void *)&v1 );
    pthread_create( &t2, NULL, thread_function, (void *)&v2 );
    pthread_join( t1, NULL );
    pthread_join( t2, NULL );
    printf( "All threads returned\n" );
    return( 0 );
}
```

Putting it all together ...

```
typedef struct {
    int num;
    const char *str;
} MY_STRUCT;


void * thread_function( void * arg ) {
    MY_STRUCT *val = (MY_STRUCT *)arg; // Cast to expected type
    printf( "Thread %lx has vaules %x,%s\n", pthread_self(), val->num, val->str );
    pthread_exit( &val->num );
}

int main( void ) {
    MY_STRUCT v1 = { 0x12345, "Val 1" };
    MY_STRUCT v2 = { 0x54312, "Val 2" };
    pthread_t t1, t2;
    printf( "Starting threads\n" );
    pthread_create( &t1, NULL, thread_function, (void *)&v1 );
    pthread_create( &t2, NULL, thread_function, (void *)&v2 );
    pthread_join( t1, NULL );
    pthread_join( t2, NULL );
    printf( "All threads returned\n" );
    return( 0 );
}
```

```
$ ./concurrency
Starting threads
Thread 7f51c3e05700 has vaules 54312,Val 2]
Thread 7f51c4606700 has vaules 12345,Val 1]
All threads returned
$
```

Thread pooling

- A systems design in which a number of threads are pre-created and assigned to requests/tasks as needed
 - Once they are done processing, they go back on the queue



Idea: amortize the cost of thread creation over many requests

Challenge ...

- Consider the following function:
- Now suppose thread “a” executes the following
- ... at the same time thread “b” executes
- What is the final value of `accounts[1]`?

Challenge ...

- Consider the following function:

```
int accounts[5] = { 0, 0, 0, 0, 0 };  
int addvalue( int account, int amount ) {  
    int num = accounts[account];  
    num = num+amount;  
    accounts[account] = num;  
    return( 0 );  
}
```

- Now suppose thread “a” executes the following
- ... at the same time thread “b” executes
- What is the final value of `accounts[1]`?

Challenge ...

- Consider the following function:

```
int accounts[5] = { 0, 0, 0, 0, 0 };  
int addvalue( int account, int amount ) {  
    int num = accounts[account];  
    num = num+amount;  
    accounts[account] = num;  
    return( 0 );  
}
```

- Now suppose thread “a” executes the following

```
addValue( 1, 100 );
```

- ... at the same time thread “b” executes

- What is the final value of `accounts[1]`?

Challenge ...

- Consider the following function:

```
int accounts[5] = { 0, 0, 0, 0, 0 };  
int addvalue( int account, int amount ) {  
    int num = accounts[account];  
    num = num+amount;  
    accounts[account] = num;  
    return( 0 );  
}
```

- Now suppose thread “a” executes the following

```
addValue( 1, 100 );
```

- ... at the same time thread “b” executes

```
addValue( 1, 200 );
```

- What is the final value of `accounts[1]`?

Watch the execution!

Thread A	Thread B	num "A"	num "B"	accounts[1]
int num = accounts[account];		0		0
	int num = accounts[account];	0		0
	num = num+amount;		200	
	accounts[account] = num;		200	200
num = num+amount;		100		
accounts[account] = num;		100		100

- Q:What happened?
- A:The temporary value of shared held stomped on the other thread's data/computation.
 - This is a known as a race condition
- The OS course will teach you how to handle this using synchronization via MUTEXes (mutual exclusion)

Thread tradeoffs

- Benefits
 - ▶ straight-line code, line processes or sequential
 - ▶ still the case that much of the code is identical!
 - ▶ parallel execution; good CPU, network utilization
 - ▶ lower overhead than processes
 - ▶ shared-memory communication is possible
- Disadvantages
 - ▶ synchronization is complicated
 - ▶ shared fate within a process; one rogue thread can hurt you
 - ▶ security (no isolation)